

UNIVERSITÉ LIBRE DE BRUXELLES
Faculté des Sciences
Département d'Informatique

Protein Residue Contact Prediction based
on a Fully-Convolutional Neural Architecture

Antoine Passemiers

Promotor : Prof. Tom Lenaerts
Co-supervisor: Dr. Daniele Raimondi

Master Thesis in Computer Sciences

“Our cells engage in protein production, and many of those proteins are enzymes responsible for the chemistry of life.”

Randy Schekman, 2013

“To deal with hyper-planes in a 14-dimensional space, visualize a 3-D space and say “fourteen” to yourself very loudly. Everyone does it.”

Geoffrey Hinton, 2014

Acknowledgment

I would first like to thank my thesis supervisor Prof. Tom Lenaerts from the Computer Science Department of Université Libre de Bruxelles (ULB), the Machine Learning Group and the Interuniversity Institute of Bioinformatics in Brussels (*IB*²). He consistently allowed me to work on everything that interested me, while still driving me back to the subject under discussion in this final version of the thesis.

I would also like to thank my co-supervisor Dr. Daniele Raimondi from the *IB*² for all the valuable advice he gave me over the past two years of master thesis preparation and writing. His expertise in structural biology and feedback gave me a lot to think about, as well as the opportunity to come up with a rigorous methodology.

I would like to thank the Machine Learning Group and *IB*² and in particular Prof. Gianluca Bontempi, as well as Raphaël Leplae from the computing centre of ULB for giving me access to Hydra (the HPC cluster of VUB/ULB), without which I would never have been able to produce the results presented in this document.

Finally, I would also like to acknowledge my classmate Robin Petit (merci Robin) for his great support and for the time he invested in proofreading the thesis when I was halfway through it.

Antoine Passemiers

Contents

1	Introduction	1
1.1	Context and objectives of the thesis	1
1.2	Contributions	3
2	Background	5
2.1	Protein contact maps	5
2.1.1	Definition	5
2.1.2	An alternative representation: protein contact networks	6
2.2	Deep learning	7
2.2.1	A definition of deep learning	7
2.2.2	The backpropagation algorithm	8
2.2.3	Fully-connected layers	10
2.2.4	Convolutional layers	11
2.2.5	Activation functions	12
2.2.6	Batch normalization	13
2.2.7	Regularization	14
2.2.8	Optimization algorithms	16
3	State-of-the-art	18
3.1	Direct Coupling Analysis	18
3.1.1	Potts model	18
3.1.2	Exact inference is hard	19
3.1.3	Pseudo-Likelihood Maximization	19
3.1.4	Gaussian Direct Coupling Analysis	20
3.2	PSICOV	23

3.2.1	Gaussian graphical models	23
3.2.2	Evolutionary couplings	23
3.2.3	Inference	24
3.3	Deep learning in Protein Contact Prediction	27
3.3.1	Input features	27
3.3.2	Features for the proposed approach	30
3.3.3	Recurrent networks	32
3.3.4	Fully-convolutional networks	32
3.3.5	Residual Networks (ResNets)	32
3.3.6	Deep fully-convolutional residual networks	33
3.3.7	U-net architecture	33
3.3.8	Dense networks (DenseNets)	34
3.3.9	TiramiProt	35
3.3.10	DeepConPred2	35
3.3.11	Properties of DL approaches	35
3.4	Contact-assisted protein folding	37
4	Materials and methods	38
4.1	Datasets	38
4.1.1	PSICOV Dataset	38
4.1.2	PConsC3's benchmark set	39
4.1.3	PConsC4's training set	39
4.1.4	Validation set	39
4.1.5	Test sets	39
4.1.6	Homology reduction	39
4.1.7	Feature extraction	41
4.2	Summary of input features	42
4.3	Proposed architecture	43
4.4	Evaluation	44
4.5	Hyper-parameter optimization	45
4.5.1	Tree-structured Parzen estimators	45
4.5.2	Search space	46
4.6	Implementation	47

4.6.1	Availability	47
5	Results	49
5.1	Hyper-Parameter Optimization	49
5.2	Model evaluation on the different benchmark sets	52
5.2.1	CASP11 targets	52
5.2.2	CAMEO targets	53
5.2.3	Membrane proteins	53
5.3	Sensitivity to the number of sequence homologs	54
5.3.1	CASP11 targets	54
5.3.2	CAMEO targets	55
5.3.3	Membrane proteins	56
5.4	Further analysis on hard CAMEO targets	57
	Appendices	62
A	Structure modelling	63
A.1	Protein 3D alignment and evaluation metrics	63
A.2	Contact-assisted 3D modelling	64
A.2.1	Graph distances	64
A.2.2	Approximate Euclidean distances	64
A.2.3	Gaussian restraints	64
A.2.4	Optimization algorithm	65
A.2.5	Evolutionary algorithm	65
A.3	Evaluation of GDE-GaussFold	66

Chapter 1

Introduction

1.1 Context and objectives of the thesis

Proteins are large macromolecules in the form of chains of building blocks called amino acids, linked by peptide bonds. Because peptide bonds are formed through a dehydration reaction between the carboxyle group of an amino acid and the amino group of another amino acid, the resulting amino acids have no free group left (except the side chain) and are called residues. Because this thesis focuses on proteins and not amino acids alone, the terms "amino acid" and "residue" will be used interchangeably.

There are 20 common amino acid types, but certain proteins may contain 2 additional amino acid types, namely pyrrolysine and selenocystein. Proteins are responsible for a wide range of functions within living organisms, including enzyme catalysis, transport of molecules, DNA replication, DNA repair, DNA transcription or cell signaling. Over 5000 types of biochemical reactions have been shown as being catalyzed by enzymes [82], which are mostly proteins. The number a ligands a protein can bind, namely the enzymatic specificity, can be determined by the structure of the protein itself [75]. The protein region binding a substrate and containing the residues involved in the catalytic process is called the active site. Enzyme structure is thus of great importance since enzymatic specificity is crucial in novel drug discovery: molecules present in tested drugs are expected to have a structure with as large as possible specificity in order to avoid unwanted effects on the patient.

According to Anfinsen's dogma [5], the structure of a protein is uniquely determined by its underlying amino acid sequence, at least when observed in protein's native environment. When moved from an unfavourable environment (where proper folding conditions are not met) to a solvent at neutral pH, a random coil (a sequence of amino acid residues oriented in random directions) will evolve towards the three-dimensional structure that minimizes Gibbs free energy. This process is called protein folding and has, however, a few exceptions.

Protein structure is organized hierarchically: primary structure, secondary structure, tertiary structure and quaternary structure. Primary structure refers to the chemical composition of the protein, hence the sequence of amino acids present in it. Secondary structure relates to the presence of structures that are local to the amino acids them-

selves: these structures are mostly α -helices and β -sheets, but multiple other structural classes exist. Tertiary structure contains information about the three-dimensional structure of the protein and results from interactions between side chains of some pairs of amino acids, such as hydrogen bonds, ionic bonds or disulfide bridges. Quaternary structure is specific to proteins having multiple polypeptide chains and relates to the structure due to intermolecular interactions between these chains. Protein contact prediction (PCP) helps predicting the tertiary structure as three-dimensional models can be reconstructed from protein contact maps (see section 3.4). Protein contact maps are a more simplistic and robust description of a protein’s geometry because they are invariant to rotations and translations. This simplification helps making deep learning methods perform well on structure prediction.

Assuming the protein backbone has no structural restriction and is composed of n residues held together by $n - 1$ peptide bonds, then the protein has $2(n - 1)$ bond angles that can be each in three different stable states. Therefore, there are at most $3^{2(n-1)}$ possible configurations, and it would take the age of the universe to find the correct folding by enumeration. There is strong evidence that protein folding is a NP-hard problem [39]. However, in practice small proteins are able to fold into a stable conformation in a fraction of a millisecond. This observation is known as the Levinthal’s paradox. There has been a long standing perspective that protein folding is guided by heuristics composed of local interactions [58]. Heuristic folding leads to misfolded proteins that can potentially cause genetic diseases. Luckily, some proteins are assisted by molecular chaperones during their folding process [27] to attain their functional conformation. It must be noted that Anfinsen’s observations of polypeptide chains refolding spontaneously in an aqueous medium have been made in the framework of in vitro studies: they do not take into account protein-protein interactions and thus cannot generalize the self-assembly process well.

Computational methods are important in structural biology, as they help in assigning biochemical or biological functions to proteins in an automated manner. The three-dimensional structure of a protein is more conserved than the underlying amino acid sequence across evolution, partly due to the fact that mutations between amino acids having the same physico-chemical properties are frequent. Prompted by this knowledge, similar functions can be assigned to proteins with low structural dissimilarity. Precisely identifying the role played by each protein in an organism is the first step towards understanding complex body mechanisms like muscle contraction, digestion or perceiving light. Also, determining the static structure of proteins help in detecting misfolded proteins which are possibly involved in diseases like Parkinson’s or Alzheimer’s, but also in diagnosing those diseases [31]. Finally, solving the protein folding problem will enable better protein design, for example to engineer enzymes like PETase so they have faster plastic-degrading capabilities [29].

Protein Contact Prediction can help determining the three-dimensional structure of proteins by limiting the search space to certain conformations that are constrained by the predicted contact maps: this methodology is called contact-assisted protein folding. The problem of predicting the structure of a protein can start by a PCP stage because the latter is a much simpler problem (despite in itself being very hard), and only a few correctly predicted contacts are sufficient to reconstruct the whole structure [50]. There are multiple well-established pipelines for the structural prediction of a newly observed protein, such as RaptorX server [74].

Most state-of-the-art PCP methods can be roughly divided into two categories: the ones based on Evolutionary Coupling Analysis (ECA) and the ones that infer contacts using supervised machine learning. In the former case, amino acid pairwise mutations are statistically modelled and the underlying model’s parameters are optimized through log-likelihood maximization or the optimization of any similar metric. In the second case, deep neural architectures are used to refine predictions made by low-level predictors such as ECA, in order to generate high-quality contact maps.

Ultimately, PCP should help making *ab initio* structure prediction in a single pass. However, most recent methods rely on a whole raft of database search, alignment, prediction and folding tools. Given a protein encoded in FASTA format, ECA is only possible using a Multiple Sequence Alignment (MSA) of this target protein against homologous proteins. These homologous proteins usually come from the same protein family as the target protein. This can be done by matching the target sequence to a Hidden Markov Model (HMM) profile representing a family like in Pfam database [26]. Once the homologous sequences have been retrieved, they have to be aligned to the target sequence using an MSA tool like HHblits or HMMER. In the next step, evolutionary couplings are extracted from the MSA using an ECA predictor like PSICOV [47] or plmDCA [24]. Eventually, predictions are gathered and refined using a deep neural architecture, necessitating the use of a deep learning framework. These successive layers of dependencies are not making PCP a straightforward process. Therefore, it seems to be a natural choice to set as an objective for this thesis the development of a predictor with minimal requirements and performance close to state-of-the-art techniques.

1.2 Contributions

During the writing of this thesis, I’ve been confronted with the need to adopt a full workflow for data retrieval and pre-processing and to develop a supervised model for accurate protein contact prediction. In order to be able to compete with state-of-the-art models, I had recourse to deep residual neural networks [40] and implemented them in a fully-convolutional manner. Indeed, the local context of a residue pair can be captured by stacking many convolutional layers and therefore increasing the receptive field [44] of the output layers of the network. The number of neighbouring input values made visible to a same hidden neuron, or receptive field, increases linearly with the depth of the neural network when using standard convolutional filters. Therefore, in order to reach the context size required to capture the long-range information of large proteins, a large number of layers had to be considered. With the aim of overcoming common issues encountered when growing very large architectures, residual connections [40] (legitimizing the use of a ResNet) as well as batch normalization [45] have been introduced. Batch normalization helps preventing internal covariate shift, a shift in the distribution of a layer’s output due to the update of the parameters, which causes the inability of the next layers to learn efficiently.

Finally, in order to promote and facilitate academic research on the topic, I have in all modesty open-sourced all the work I’ve done during the year of thesis writing. Due to my computer science background, it is my belief that the available code could serve as source of inspiration for biologists and bioinformaticians with lower capabilities in programming.

1.3 Structure of the thesis

Let's describe the global view of the thesis itself. Firstly, common state-of-the-art ECA techniques will be described, such as Direct Coupling Analysis (DCA) and Pseudo-Inverse Covariance matrices (PSICOV) (both are statistical methods based on graphical models), as well as the backpropagation algorithm and deep learning concepts involved in the design of the model developed during the thesis. In order to gain a deeper insight on best-performing methods, more details will be given for some specific deep learning methods, such as their architecture, input features and preprocessing. The generic architecture in use for this thesis will be detailed, as well as the hyper-parameter optimization procedure used for cross-validation. Finally, results will be presented in multiple sections:

- Since the proposed deep learning approach relies on DCA predictions as input features, the first step should demonstrate that deep learning is capable of refining contact maps by looking at complex visual patterns that cannot be captured by linear models.
- It should be brought to light whether deep learning's performance is less sensitive to the effective number of homologous sequences than DCA methods. The notion of effective number of homologous sequences will be introduced in section 3.3.1.1.2.
- Finally, a benchmark will be established to assess the performance of the proposed method in comparison with other supervised approaches, including state-of-the-art deep learning architectures.

Chapter 2

Background

2.1 Protein contact maps

2.1.1 Definition

A **contact** between two residues occurs when two amino acid residues from a same protein are separated by a distance below a given threshold. Let C_α be the first carbon atom attached to a functional group and C_β is the first carbon atom attached to C_α . Functional groups of amino acid residues can be either amine ($-NH_2$) or carboxyl ($-COOH$). In the first case, the amino acid is called alpha amino acid and has an amine group directly attached to the C_α of the carboxyl group. In the second case, it is called beta amino acid and has an amine group attached to the C_β of the carboxyl group.

The distance metric can be either the distance between $C_\alpha - C_\alpha$ atoms or the distance between $C_\beta - C_\beta$ atoms. It should be underlined that glycine does not have a C_β and thus C_α is used instead.

Most common distance thresholds range between 6 and 12 Å. An angstrom (Å) is an unit of length equivalent to 10^{-10} m, or 10^{-1} nm. Therefore the notion of residue contact depends to a large extent on the threshold used. For example, the average percentage of contacts in the 150 proteins reported in the original PSICOV article [47] is equal to 7%, 14%, 26%, 39% with thresholds 7, 10, 13 and 16 Å, respectively.

The present thesis complies with the definition of contact maps as given by the Critical Assessment of methods of protein Structure Prediction (CASP) [30]: two residues are in contact if their C_β (C_α for glycine) atoms are separated by a distance below a threshold of 8 Å.

By extension, a **protein contact map** can be defined as a symmetric binary matrix C where element C_{ij} is equal to 1 if residues i and j are separated by a distance below the given threshold, and 0 otherwise. Contact maps are invariant to rotations and easier to predict with machine learning methods, contrary to matrices of pairwise distances. Furthermore, the original 3D residue coordinates can be recovered from contact maps [95] as explained in section 3.4. Such hypothesis is in line with Anfinsen's Dogma which postulates that the secondary and tertiary structures of a protein can be inferred

from its primary structure: even after disrupting the hydrophobic bonds of a protein, experiments suggest that the latter can recover its original structure with some assisted folding, highlighting the idea that tertiary structure is encoded in the sequence of amino acids itself. The task of predicting contact maps is illustrated in figure 2.1.




Figure 2.1: (Left) Ground-truth contact map of PDB:1CXYYA defined at a contact threshold of 8 Å. (Center) PConsC2 [89]’s predicted contact map for the same protein. (Right) Evaluation of the predicted contact map. True positives are colored in green and false positives are colored in red. Only top L residue pairs are considered in the evaluation, where L is the protein length.

Left image is a ground-truth contact map associated to the 3D coordinates of a protein, obtained either by X-ray crystallography or nuclear magnetic resonance (NMR). Middle image is a matrix of predicted contact probabilities and right image corresponds to the evaluation of the predicted contact map on the basis of the ground-truth contact map.

2.1.2 An alternative representation: protein contact networks

Another way of interpreting contact maps is viewing them as adjacency maps of protein contact networks. This sub-section will present the formal definition of protein contact networks as suggested in [19]. Let $G = (V, E)$ be a graph where V is a set of vertices and E a set of edges. Such a graph G can be encoded as a matrix A called the adjacency matrix. Given a set of vertices $\{v_1, \dots, v_n\}$, adjacency matrix $A \in \{0, 1\}^{n \times n}$ is such that:

$$A_{ij} = \begin{cases} 1 & \text{if } (v_i, v_j) \in E \\ 0 & \text{otherwise} \end{cases} \quad (2.1)$$

Using this definition, many adjacency matrices of a same graph exist. Indeed, many matrices can be created by simply making permutations of rows and columns. However, the ordering of vertices is determined by the sequence of amino acids, making it unique. Weighted graphs are slightly different than regular graphs since they are defined not only by their connections but also their weights. Accordingly, the adjacency matrix of a weighted graph is adapted as follows:

$$A_{ij} = \begin{cases} w_{ij} & \text{if } (v_i, v_j) \in E \\ 0 & \text{otherwise} \end{cases} \quad (2.2)$$

where w_{ij} is the weight of edge (v_i, v_j) .

Also, the degree $\deg(v_i)$ of a vertex v_i is defined as the number of neighbouring vertices, or in other words the number of vertices each sharing an edge with v_i :

$$\deg(v_i) = \sum_{j=1}^n A_{ij} \quad (2.3)$$

This definition of vertex degree also holds for weighted graphs. However, many authors favor minimal representation of protein structure and abandon the use of weights. The diagonal degree matrix D can be defined by the following relation:

$$D_{ij} = \begin{cases} \deg(v_i) & \text{if } i = j \\ 0 & \text{otherwise} \end{cases} \quad (2.4)$$

A **protein contact network** is a graph having its set of vertices ordered by primary structure, where each vertex is an amino acid itself. The presence of an edge between two such vertices indicates that the two corresponding residues are in contact. Such a network is useful to make a compact representation of a protein structure and metrics such as path length or graph diameter are important for the analysis of long-range residue interactions [19].

Let sp_{v_1, v_2} be the number of vertices located on the shortest path from v_1 to v_2 , called the distance between v_1 and v_2 . The diameter of a graph $G = (V, E)$ is defined as follows:

$$\text{diam}(G) = \max\{sp_{v_1, v_2} | v_1, v_2 \in V\} \quad (2.5)$$

This protein contact network formalism has been used in the design of GDFuzz3D [76] for contact-assisted protein folding, which has shown remarkable performance on the PSICOV dataset [47]. The notion of shortest path length has been used as a measure of real euclidean distance, as will be described in section 3.4.

2.2 Deep learning

2.2.1 A definition of deep learning

Beyond the trendy words, it is quite difficult to find a consensus on the definition of deep learning. The concept is often associated to the concept of inferring a high-level representation of the data by alternating many times between parameterized functions and non-linearities. Deep artificial neural networks serve this purpose well since they are composed of many sets of parameters and a large stack (or graph) of mathematical

operators linked to an objective function to be optimized. Each parameteric operator may rely on a subset of the network parameters.

In most simple cases (e.g. feedforward neural network), the network can be described as a regular stack of operators. As a result, the objective function is a composition of all the underlying mathematical operations. Such a network is usually trained using the backpropagation algorithm. The latter method consists in minimizing the objective function, which usually is a dissimilarity measure between what the network predicts for a given input and what the human supervisor expects for such input. More specifically, backpropagation is an iterative algorithm that evaluates the gradient of the objective at each iteration and performs one step in the direction of the steepest descent in the parameter space. The algorithm is expected to stop once a global minimum has been reached. Formal details about the algorithm are going developped in the next sub-section.

Deep learning is also often viewed as the ability of a machine to build a hierarchical representation of the data by mapping input values to high-level features. According to Yoshua Bengio and Yann LeCun, neural networks only exemplify the notion of deep architectures. They provided a sufficiently good basis for a definition:

Deep architectures are compositions of many layers of adaptive non-linear components, in other words, they are cascades of parameterized non-linear modules that contain trainable parameters at all levels. Deep architectures allow the representation of wide families of functions in a more compact form than shallow architectures, because they can trade space for time (or breadth for depth) while making the time-space product smaller, as discussed below. The outputs of the intermediate layers are akin to intermediate results on the way to computing the final output. Features produced by the lower layers represent lower-level abstractions, that are combined to form high-level features at the next layer, representing higher-level abstractions [9].

This definition seems to be perfectly appropriate for neural networks since they are precisely made of linear - and consequently parametric - operations followed by activation functions which are non-linear by nature.

2.2.2 The backpropagation algorithm

In this sub-section, the backpropagation algorithm is going to be introduced formally in order to understand the subtleties of deep learning. Explanations will focus w.l.o.g. on feedforward networks and be accompanied by my own formalism. For more details on the backpropagation algorithm, the reader can refer to a more general and theoretical description from Yann LeCun [56] or a more pictorial version [54].

Let's consider a neural network containing no cycle. Each of its layers can be viewed as a couple $(f_i(\theta_i, X), b_i(\theta_i, S(X)))$, where f_i is the forward pass function of layer i used for predicting, b_i is the backward pass function, θ_i is the set of parameters, and X, Y are input tensors of shapes compatible with f_i and g_i , respectively, and $S(X)$ is the signal tensor propagated from next layer back to current layer. Let's make the assumption that convolutional layers are two-dimensional and that input instances are

image-like data. (one-dimensional and three-dimensional convolutions can be described analogously). Also, let's consider a particular case of neural network (called feedforward) consisting of a stack of neural layers with no feedback connections: let's note that the framework can still be easily extended to the recurrent case. Let b be the number of examples in the input tensor (more commonly referred to as the batch size), w and h respectively the width and height of the images, and c the number of channels. Finally, let n be the number of layers and m be the number of output neurons in the network. Knowing this, the output $Y \in \mathbb{R}^{b \times m}$ of the network can now be described as such:

$$Y = (\bigcirc_{i=1}^n f_{\theta_i})(X) \quad (2.6)$$

where $X \in \mathbb{R}^{b \times w \times h \times c}$ and $f_{\theta_i}(X)$ is syntactic sugar for denoting $f_i(\theta_i, X)$ in a more convenient way. It can be observed that the prediction function of the network is basically a large composition of functions.

Such model is designed to optimize a function reflecting its ability to accurately predict a target value or to abstractly represent the input data in a more general sense. Accordingly, let's introduce a generic loss function $L(Y) : \mathbb{R}^{b \times m} \rightarrow \mathbb{R}$ that measures the model's inability to fulfill the given task. The loss takes the output Y of the network as input, and represents the objective function to be minimized. Using the composition rule and by replacing Y in $L(Y)$, we obtain the following expression:

$$\hat{\theta} = \underset{\theta \in \Theta}{\operatorname{argmin}} L((\bigcirc_{i=1}^n f_{i, \theta_i})(X)) \quad (2.7)$$

where Θ is the set of all possible values for the parameter set $\theta = (\theta_1, \dots, \theta_n)$. The generic task of minimizing a scalar continuous function can be achieved using numerous continuous optimization techniques among gradient descent algorithms [80] or quasi-Newton methods [12], as will be detailed in section 2.2.8. In practice, gradient descent approaches require more iterations to converge to a satisfying solution, but are easier to implement. Also, contrary to quasi-Newton methods, they don't require to implicitly compute the hessian matrix of the loss function according to the network parameters, which makes them less computation-intensive.

Let's consider the optimization of the loss function in the gradient descent framework. The loss function is minimized by moving in the parameter space in the direction of the loss gradient, with a step proportional to the learning rate (a parameter either determined empirically or adjusted dynamically during optimization phase). Luckily, since we are regarding our neural network as a stack of layers (viewed as a composition of functions), the gradient computation can be decomposed using the chain rule:

$$\frac{\partial(f \circ g)}{\partial w}(X) = \nabla f(g(X)) \cdot \frac{\partial g}{\partial w}(X) \quad (2.8)$$

where w can be any parameter of the network. Knowing this, the gradient of the loss function w.r.t. to the parameter set θ_p of layer j (for any layer j with learnable parameters), can be decomposed as the following product:

$$\prod_{k=1}^p f'_{\theta_k} \left((\odot_{i=1}^k f_{\theta_i})(X) \right) \cdot L' \left((\odot_{j=1}^n f_{\theta_j})(X) \right) \quad (2.9)$$

Each factor k of the product can be computed using the definition of function f'_k , and the current input to layer k . However, layer k requires the factor from layer $k+1$ in order to compute loss gradient according to its own parameters. Consequently, the signal (the product of factors accumulated from layer n to current layer i) is passed from layer $i+1$ to layer i . In a more general sense, the gradient signal is passed from the output layer to the input layer, hence the name "backpropagation".

The move in the gradient direction with step α (the so-called learning rate) is such that:

$$\theta_k \leftarrow \theta_k - \alpha \cdot \nabla_{\theta_k} L(X) \quad \forall k \in \{1, \dots, n\} \quad (2.10)$$

This step is repeated until one of the stop criteria has been met. For example, the algorithm stops when a maximum number of iterations has been reached. However, gradient descent is not the only optimization algorithm that yields satisfying results in practice. For example, Limited-memory BFGS [12] relies on a second order approximation of the loss function given a limited number of past update vectors: this provides a better search direction but in return does not theoretically guarantee that the loss function actually decreases at each iteration.

2.2.3 Fully-connected layers

A Multi-layer perceptron is a neural network composed of multiple layers, where each layer's forward pass consists of a linear combination of the inputs followed by an element-wise non-linear activation function. Let $X^{(p)} \in \mathbb{R}^{n \times m}$ be the input matrix of layer p , $W \in \mathbb{R}^{m \times k}$ the weight matrix, $b \in \mathbb{R}^k$ the bias vector, $n^{(p)}$ the number of examples fed as input to layer p and σ the non-linear activation function of layer p . Each layer can than be formalized as follows:

$$X_{i,k}^{(p+1)} = \sigma \left(\sum_{j=1}^{n^{(p)}} X_{i,j}^{(p)} W_{j,k} + b_k \right) \quad (2.11)$$

Backpropagation requires to compute the partial derivatives of layer outputs with respect to current layer parameters:

$$\frac{\partial X_{i,k}^{(p+1)}}{\partial W_{j,k}} = \sigma' \left(\sum_{j=1}^{n^{(p)}} X_{i,j}^{(p)} W_{j,k} + b_k \right) X_{i,j}^{(p)} \quad (2.12)$$

$$\frac{\partial X_{i,k}^{(p+1)}}{\partial b_k} = \sigma' \left(\sum_{j=1}^{n^{(p)}} X_{i,j}^{(p)} W_{j,k} + b_k \right) \quad (2.13)$$

where $\sigma'(x)$ is the derivative of $\sigma(x)$, typically $\sigma(x)(1 - \sigma(x))$ for the sigmoid function.

Multi-layer perceptrons have been proved to be Universal Approximators [42], meaning that they can approximate feedforward prediction functions that minimize any training loss (loss function computed on the training set). However, this fact does not inform about the type of non-linear function to use in order to minimize a given loss function. More importantly, this does not guarantee that the model will perform well on unseen examples. Indeed, high representational power is required when the classification task is abstract. To overcome this problem and lower the validation loss as much as possible, data scientists usually stack more layers on top of each other, but this may imply high computational requirements. Convolutional layers are used instead of dense weight matrices.

2.2.4 Convolutional layers

One of the major advances in semantic segmentation is due to Convolutional Neural Networks (CNNs) [33]. A CNN is an artificial neural network made of a stack of neural layers [55]. One characteristic of CNNs is the presence of convolutional filters that map raw data to more abstract features. Each filter (or kernel) is locally connected to its output unit, which allows the convolutional layer to capture some local information about the inputs, as opposed to fully-connected layers that don't take any spatial information into account when passing data forward. This procedure is inspired by the notion of **receptive field** introduced by Hubel and Wiesel [44]: by stacking multiple convolutional layers, the number of values (e.g. residue pair features) made visible to a same hidden neuron, namely the receptive field, increases with the number of layers. The property of being able to increase the receptive field with the depth of the network is convenient when a high representational power is required, which is the case for the task of protein contact prediction. Indeed, a receptive field equal to the protein size is needed in order for the model to take into account the global protein structure.

Weights are no longer stored in a bidimensional matrix since all inputs are no longer connected to each neuron of the current layer. Instead, each neuron is connected to a certain neighborhood of inputs. In this way, the network drastically reduces its number of parameters but still takes the spatial dependence of the data into account. If the convolutional layer is designed for processing multi-channel images for example, the parameters will be stored in a 4-dimensional tensor. Let $W \in \mathbb{R}^{b \times h \times w \times n_c}$ be the weights of the convolutional filters, $X^{(p)} \in \mathbb{R}^{b \times h_b \times w_b \times n_c}$ the input images of layer p , $b \in \mathbb{R}$ the bias vector, and $X^{(p+1)} \in \mathbb{R}^{b \times (\lfloor (h_b - h)/\beta_1 \rfloor + 1) \times (\lfloor (w_b - w)/\beta_2 \rfloor + 1) \times n_f}$ the output feature maps. n_c is the number of channels, h is the filter height, w is the filter width and (β_1, β_2) are the strides (vertical and horizontal distances between neighboring pixels in the neighborhood connected to a same neuron). Let's consider the relation between the input images and the output feature maps:

$$X_{i,j,k,l}^{(p+1)} = \sigma \left(\sum_{\alpha=1}^h \sum_{\delta=1}^w \sum_{c=1}^{n_c} W_{j,\alpha,\delta,c} X_{i,k+\beta_1\alpha,l+\beta_2\delta,c}^{(p)} + b_j \right) \quad (2.14)$$

where i is the image identifier and j is the filter index. Partial derivatives are given by:

$$\frac{\partial X_{i,j,k,l}^{(p+1)}}{\partial W_{j,\alpha,\delta,c}} = \sigma' \left(\sum_{\alpha'=1}^h \sum_{\delta=1}^w \sum_{c=1}^{n_c} W_{j,\alpha',\delta,c} X_{i,k+\beta_1\alpha',l+\beta_2\delta,c}^{(p)} + b_j \right) X_{i,k+\beta_1\alpha',l+\beta_2\delta,c}^{(p)} \quad (2.15)$$

$$\frac{\partial X_{i,j,k,l}^{(p+1)}}{\partial b_j} = \sigma' \left(\sum_{\alpha'=1}^h \sum_{\delta=1}^w \sum_{c=1}^{n_c} W_{j,\alpha',\delta,c} X_{i,k+\beta_1\alpha',l+\beta_2\delta,c}^{(p)} + b_j \right) \quad (2.16)$$

Just as in the case of fully-connected layers, the computations for the signal propagation are not shown because this report is intended to remain brief.

2.2.5 Activation functions

An activation function describes the output value of a neuron and is biologically inspired. It is a mathematical representation of the level of action potential sent along its axon. More formally, it is a non-linear scalar function that takes a scalar as input. The presence of activation functions in neural networks along with fully-connected layers allows them to increase their representational power. Indeed, a stack of fully-connected layers without activation functions would have the same representational power as a single fully-connected layer, since a linear combination of linear combinations is itself a linear combination. Thus, activation functions help to actually build a hierarchical representation of the data by curving the projected space multiple times and at each layer.

However, not every activation function is suitable for backpropagation and one of the reasons for the success of deep learning is the low computational requirements for the gradients. Most of the activation functions are non-parametric and element-wise, which makes it easy to compute the signal during backward pass.

The best known activation function is the sigmoid function $\sigma(x)$. It has the property to have a derivative $\sigma'(x)$ expressed as a function of $\sigma(x)$, which speeds up computation times, assuming that the neural outputs are cached.

$$\begin{aligned} \sigma(x) &= \frac{1}{1 + \exp(-x)} = \frac{\exp(x)}{1 + \exp(x)} \\ \sigma'(x) &= \sigma(x)(1 - \sigma(x)) \end{aligned} \quad (2.17)$$

However, LeCun [57] does not recommend standard sigmoid functions because normalizing activation functions generally ensure better performance. For this reason, the hyperbolic tangent is suitable because its outputs are centered around zero. Also, its derivative $\tanh'(x)$ is expressed as a function of $\tanh(x)$ which is computationally convenient. Finally, an additional linear term can be added in order to avoid flat areas, leading to an activation function of the following form: $f(x) = \tanh(x) + ax$.

$$\begin{aligned} \tanh(x) &= \frac{\exp(x) - \exp(-x)}{\exp(x) + \exp(-x)} = \frac{\exp(2x) - 1}{\exp(2x) + 1} \\ \tanh'(x) &= 1 - \tanh^2(x) \end{aligned} \quad (2.18)$$

Assuming that target values are in the set $\{-1, 1\}$ in the framework of binary classification, the hyperbolic tangent can be linearly modified to obtain a new function of the form $f(x) = 1.7159 \tanh(\frac{2}{3}x)$. Such an activation function is profitable because, has its second derivative maximized at $x = -1$ and $x = 1$, avoiding saturation effects.

The chain rule informs us that the gradient of a given layer is factorized as a product of vectors/matrices computed by next layers. Because the absolute values of a layer's outputs are always less than one for both tanh and standard sigmoid activation functions, but also the absolute values of the gradient's components, deep architectures are often subject to vanishing gradients. Linear rectifier units (ReLU) are piecewise linear functions designed to solve these issues by keeping positive inputs unchanged. Let's note that ReLU is not differentiable at $x = 0$ but inputs can be reasonably assumed to be rarely equal to zero in practice.

$$\begin{aligned} \text{ReLU}(x) &= \max(x, 0) \\ \text{ReLU}'(x) &= \begin{cases} 1 & \text{if } x > 0 \\ 0 & \text{if } x < 0 \end{cases} \end{aligned} \quad (2.19)$$

The outputs of a neural network are often desired to sum to one, especially when the classification task is to assign each class to a probability conditionally to the network's input. In the case where there are m classes, the output layer is composed of m neurons where the activation function associated to neuron i is given by:

$$\begin{aligned} \sigma(x_i) &= \frac{\exp(x_i)}{\sum_{k=1}^m \exp(x_k)} \\ \sigma'(x_i) &= \sigma(x_i)(1 - \sigma(x_i)) \end{aligned} \quad (2.20)$$

where x_i is the component i of the output vector. This function is identical to the Boltzmann distribution introduced in section 3.1.1.

2.2.6 Batch normalization

According to Ioffe and Szegedy [45], deep neural networks are subject to a phenomenon called **internal covariate shift**. When the learning rate is too large, the distribution of a layer's output is drastically altered, making it difficult to train the next layer since the latter is constantly adapting to the new distribution. Batch normalization helps dealing with this issue and allows us to run the optimization algorithm with less careful parameter initialization and a larger learning rate.

When the network is trained with batch learning, its parameters are updated at every batch. Therefore, the distribution of each layer's output is changed at each batch. This is the reason for using the statistics of each batch individually to normalize the data between layers.

Let $\mathcal{B} = \{x_1, x_2, \dots\}$ be the input batch, x_i one of the input example and \hat{x}_i the same input example after normalization. Batch statistics are the element-wise mean vector $\mu_{\mathcal{B}}$ and element-wise standard deviation vector $\sigma_{\mathcal{B}}$ given in equations 2.21.

$$\begin{aligned}\mu_{\mathcal{B}} &= \frac{1}{|\mathcal{B}|} \sum_{i=1}^{|\mathcal{B}|} x_i \\ \sigma_{\mathcal{B}}^2 &= \frac{1}{|\mathcal{B}|} \sum_{i=1}^{|\mathcal{B}|} (x_i - \mu_{\mathcal{B}})^2\end{aligned}\tag{2.21}$$

These statistics are used for normalizing all the examples present in current batch \mathcal{B} :

$$\hat{x}_i \leftarrow \gamma \frac{x_i - \mu_{\mathcal{B}}}{\sqrt{\sigma_{\mathcal{B}}^2 + \epsilon}} + \beta\tag{2.22}$$

Variables γ and β in equation 2.22 are the scaling and location parameters used to re-scale and re-center the output distribution after normalization. The optimal values for these parameters is also find by gradient descent, in order to empirically determine the ideal mean and variance of the distribution expected by the next layer of the network. Let's note that mean γ and variance β slightly vary from one iteration to the other due to parameter update but are not directly related to batch statistics, which ensures some semblance of stationarity in the output distribution.

Partial derivatives of γ and β can be easily be found by applying the methodology used in section and using equations 2.21 and 2.22, or by referring to the original paper [45].

2.2.7 Regularization

From an optimization perspective, regularization is a penalty used to prevent parameters from growing arbitrarily big during training. According to Occam's law of parsimony, simpler hypotheses should be privileged over more complex ones. Therefore, when the neural architecture involves a large number of free parameters in the presence of relatively few data samples, regularization helps reducing parameters importance and converging to less arbitrary parameter values. From a Bayesian perspective, regularization provides a prior distribution over the model parameters. In Bayes formula, the posterior $P(\theta|X, \alpha)$ is a function of both the prior $P(\theta|\alpha)$ and the likelihood of the data $P(X|\theta, \alpha)$ under model θ .

$$P(\theta|X, \alpha) = \frac{P(X|\theta, \alpha) P(\theta|\alpha)}{P(X|\alpha)}\tag{2.23}$$

The relation between the loss function of a neural network and Bayes formula can be established by proving the two following points:

- The log-likelihood of the data is equal to the negative cross-entropy.

- The regularization term is proportional to the prior distribution of the parameters.

The first part is easy to show since negative log-likelihood can be obtained from binary cross-entropy:

$$CE(\hat{y}, y) = -\log \prod_{i=0}^n P(\hat{y}_i)^{y_i} \quad (2.24)$$

$$= -\sum_{i=0}^n y_i \log \hat{y} + (1 - y_i) \log 1 - \hat{y} \quad (2.25)$$

This allows us to provide a statistical interpretation of the loss function. Regarding priors, L_1 and L_2 regularizations are going to be introduced in the following two sections.

2.2.7.1 L_1 regularization

Adding a L_1 regularization term to the loss function reduces to providing a Laplacian prior on model parameters.

$$\max_{\theta} \log P(\theta|\eta, b) = \max_{\theta} \log \prod_{i=1}^m \frac{1}{2b_i} \exp\left(-\frac{|\theta_i - \eta_i|}{b_i}\right) \quad (2.26)$$

$$= \max_{\theta} \sum_{i=1}^m \frac{|\theta_i - \eta_i|}{b_i} - \log 2b_i \quad (2.27)$$

$$= \min_{\theta} \sum_{i=1}^m |\theta_i - \eta_i| \quad (2.28)$$

By setting vector $\eta \in \mathbb{R}^m$ to 0, the resulting regularization term takes its final well-known form $\sum_{i=1}^m |\theta_i|$.

2.2.7.2 L_2 regularization

L_2 regularization acts as a Gaussian prior on model parameters. This can be highlighted by setting the probability density function of the Gaussian distribution as the prior and show that the regularization term is proportional to the logarithm of the product of priors.

$$\max_{\theta} \log P(\theta|\eta, \sigma) = \max_{\theta} \log \prod_{i=1}^m \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{(\theta_i - \eta_i)^2}{2\sigma^2}\right) \quad (2.29)$$

$$= \max_{\theta} \sum_{i=1}^m -\frac{(\theta_i - \eta_i)^2}{2\sigma^2} - \log \sqrt{2\pi\sigma^2} \quad (2.30)$$

$$= \min_{\theta} \sum_{i=1}^m (\theta_i - \eta_i)^2 \quad (2.31)$$

Again, by setting vector $\eta \in \mathbb{R}^m$ to 0, the regularization term takes its final form $\sum_{i=1}^m \theta_i^2$.

2.2.8 Optimization algorithms

Gradient descent is a very popular optimization algorithm, but is rarely used as such in practice since state-of-the-art deep learning frameworks offer more advanced gradient-based techniques [80]. What is meant by gradient is the vector obtained by concatenation of the gradients w.r.t. each layer's parameters. This final gradient vector gives an improvement direction, but a decrease of the loss function is only guaranteed by moving by an arbitrary small step in the parameter space.

Gradient descent has three variants: batch, mini-batch and stochastic gradient descent. In batch gradient descent, all training examples are used to compute the improvement direction: this is done by computing the gradient for each training example and averaging across all examples. In mini-batch gradient descent, only a subset of training examples are being considered for the computation of the improvement direction (which can thus be seen as an approximation of the actual gradient). Usually, the ordering of training examples is shuffled at the beginning of each iteration (also called epoch) and then examples are sampled in the resulting order repeatedly, to ensure that each of them is seen by the model exactly once per iteration. In the stochastic variant, only the gradient of a single training example is used to approximate the improvement direction. Due to the high variability of gradients from one example to the other, the improvement direction is changing in a chaotic manner during the optimization process, hence the adjective stochastic. The three types of improvement vectors are summarized in table 2.1.

Name	Number of examples involved	Formula
Average (true) gradient	N	$g_t = \frac{1}{N} \sum_{i=1}^N \nabla L(f_{x_t}(Z_i))$
Mini-batch gradient	$ B $	$g_t = \frac{1}{ B } \sum_{i \in B} \nabla L(f_{x_t}(Z_i))$
Sample gradient	1	$g_t = \nabla L(f_{x_t}(Z))$

Table 2.1: Types of gradients and gradient approximations used in common optimization methods. Here the term "sample" refers to a sample of 1 example.

In its most simplistic form, gradient descent optimization consists in updating the parameter vector x_t from step t using the following rule:

$$x_{t+1} = x_t + \Delta x_t \quad (2.32)$$

$$= x_t - \eta g_t \quad (2.33)$$

where update vector Δx_t is equal to the negative approximated gradient $-g_t$ multiplied by a learning rate η . Learning rate controls how much model parameters are being updated in the improvement direction.

Stochastic gradient descent has been extended with a so-called momentum [81] term that accelerates the update when gradients approximately point in the same direction

from one step to another.

$$\Delta x_t = \rho \Delta x_{t-1} - \eta g_t \quad (2.34)$$

ρ is a decay parameter that controls how much to keep track of past update vectors. In practice, the landscape of the loss function w.r.t. parameters is likely to be composed of many narrow valleys, where standard stochastic gradient vector is most likely to be inefficient due to the small norm of gradients along valleys. The momentum helps optimizing across such valleys in a smaller number of steps due to its additive effect when gradient vectors are similar from one step to the next one.

ADADELTA [105] is an adaptive extension of stochastic gradient descent that is robust to the noise introduced by the high variability of sample gradients. Also, it dynamically selects the learning rate so that no hyper-parameter tuning is required on it.

$$\Delta x_t = -\lambda g_t = -\frac{RMS[\Delta x]_{t-1}}{RMS[g]_t} g_t \quad (2.35)$$

where λ is a dynamic learning rate that globally decays throughout the learning process. The root mean square of previously seen examples $RMS[g]_t = \sqrt{E[g^2]_t + \epsilon}$ is computed as an exponentially decaying average for memory efficiency:

$$E[g^2]_t = \rho E[g^2]_{t+1} + (1 - \rho) g_t^2 \quad (2.36)$$

where ρ is a decay parameter similar to the momentum.

Other optimization methods, like RMSProp [35], AdaGrad [21], Adam [53] or L-BFGS [12] are also efficiently implemented in most deep learning frameworks.

Chapter 3

State-of-the-art

3.1 Direct Coupling Analysis

3.1.1 Potts model

Experts have long thought that the three-dimensional structure of proteins is related to their amino acid composition. However, homologous proteins are subject to a high variability with regards their amino acid composition. Structural conservation across evolution induces stresses on these variability patterns in such a manner that spatially close residues have a restricted set of acceptable amino acid substitutions. Therefore, neighbouring residues are forced to coevolve, which results in correlated mutations [66].

Which makes structural prediction complex is that correlated mutations can be caused either by low data quality or availability manifested by the absence of a large number of homologous proteins, a high redundancy among homologous proteins or various phylogenetic effects. In addition to all these sources of error, a correlation between two residue sites may also be mediated by a third residue located at a different site and having direct correlations with the two other ones.

The core idea of Direct Coupling Analysis (DCA) is to disentangle direct correlations and indirect correlations related to residues at intermediary positions. Potts model allows to perform this disentanglement through inverse statistical mechanics [25].

DCA takes ideas from the Potts model [?], a generalization of the problem originally stated by Cyril Domb, called the Ising problem. The Potts model used in DCA involves two types of parameters:

- The pairwise **couplings** J are the values of interest for protein contact prediction since $J_{ij}(s_i, s_j)$ is defined as the predicted distance between residues i and j in the protein structure, knowing that they are in states s_i and s_j respectively [?].
- The **fields** h are local biases of the Boltzmann distribution.

Evolutionary-related sequences are modelled by the distribution given by the Maximum-Entropy Principle. Among the family of distributions that are suited for proteins (which are discrete sequences), the one that maximizes entropy is the Boltzmann distribution, with the following probability mass functiony [?]:

$$P(s|J, h) = \frac{1}{Z} \exp\left(\sum_{i=1}^L \sum_{j=i+1}^L J_{ij}(s_i, s_j) + \sum_{i=1}^L h_i(s_i)\right) \quad (3.1)$$

where couplings J and fields h are the model parameters, s is an amino acid sequence and Z is a normalization factor called partition function ensuring that the sum $\sum_s P(s|J, h)$ over all lexicographically possible sequences is equal to one. Let's note that residue s_i at site i is defined over the alphabet $\{1, \dots, q\}$, where q is the number of possible residue states (namely its amino acid type).

3.1.2 Exact inference is hard

Given a multiple sequence alignment containing M sequences, a naïve approach would be to maximize its log-likelihood:

$$\begin{aligned} \log L(J, h|s^{(1)}, \dots, s^{(M)}) &= \sum_{k=1}^M \log P(s^{(k)}|J, h) \\ &= \sum_{k=1}^M \log \left(\frac{1}{Z} \exp\left(\sum_{i=1}^L \sum_{j=i+1}^L J_{ij}(s_i^{(k)}, s_j^{(k)}) + \sum_{i=1}^L h_i(s_i^{(k)})\right) \right) \\ &= -M \log(Z) + \sum_{k=1}^M \left(\sum_{i=1}^L \sum_{j=i+1}^L J_{ij}(s_i^{(k)}, s_j^{(k)}) + \sum_{i=1}^L h_i(s_i^{(k)}) \right) \end{aligned} \quad (3.2)$$

with the following partial derivatives:

$$\begin{aligned} \frac{\partial}{\partial J_{ij}(a, b)} \log L(J, h|s^{(1)}, \dots, s^{(M)}) &= -M \frac{\partial \log(Z)}{\partial J_{ij}(a, b)} + \sum_{k=1}^M [s_i^{(k)} = a] [s_j^{(k)} = b] \\ \frac{\partial}{\partial h_i(a)} \log L(J, h|s^{(1)}, \dots, s^{(M)}) &= -M \frac{\partial \log(Z)}{\partial h_i(a)} + \sum_{k=1}^M [s_i^{(k)} = a] \end{aligned} \quad (3.3)$$

where $[\cdot]$ are Iverson brackets.

However, there is no straightforward method to compute the partition function Z or L 's gradient for large systems due to the discrete nature of amino acid sequences. Indeed, Z contains 21^L terms for systems with 21 possible symbols (amino acid types + gap) and sequences of length L . For this aim, several methods like Mean-Field (mfDCA) [68], Message Passing (mpDCA) [102], Pseudo-Likelihood Maximization (plmDCA) [24] or Multivariate Gaussian Modeling (GaussDCA) [7] have been developed.

3.1.3 Pseudo-Likelihood Maximization

plmDCA [24] addresses the problem of estimating the partition function by maximizing the pseudo-loglikelihood instead of the loglikelihood. The pseudo-loglikelihood can be

expressed as the sum of loglikelihoods $\log L(J_r, h_r)$, where each $\log L(J_r, h_r)$ is computed at a single site r . The method thus assumes the conditional independence between variables belonging to different sites. However, the partition function at a given site can be easily computed as a sum of 21 terms since a state can take 21 possible values at a given position. More formally, the penalized loglikelihood at site r is given by:

$$\begin{aligned} \log L^{(reg)}(J_r, h_r) = & -\frac{1}{M_{eff}} \sum_{k=1}^M w_k \left(h_r(s_r^{(k)}) + \sum_{i \neq k}^L J_{ri}(s_r^{(k)}, s_i^{(k)}) - \log(Z_r) \right) \\ & + \lambda_h \|h_r\|_2^2 + \lambda_J \|J_r\|_2^2 \end{aligned} \quad (3.4)$$

where
$$Z_r = \sum_{a=1}^q \exp \left(h_r(a) + \sum_{i \neq r} J_{ri}(a, s_i^{(k)}) \right)$$

It can be observed that the formula contains a L_2 penalty term for both fields and couplings, and that each log-probability is weighted by a protein weight w_k . The latter is computed as the protein contribution to set the set of effective sequences, as described in the section of M_{eff} 3.3.1.1.2. It must be observed that the optimization procedure is called asymmetric pseudolikelihood maximization because each matrix $J(i, j)$ is supposed to be symmetric and in practice estimated independently. This allows plmDCA to run in parallel by optimizing $\log L^{(reg)}(J_r, h_r)$ each on a different core.

After maximizing the pseudo-loglikelihood in parallel, all remaining information that can be explained by the fields are removed from the couplings by applying an average sum correction w.r.t. the states:

$$\hat{J}_{ij}(a, b) = J_{ij}(a, b) - \frac{1}{q} \sum_{a=1}^q J_{ij}(a, b) - \frac{1}{q} \sum_{b=1}^q J_{ij}(a, b) + \frac{1}{q^2} \sum_{a=1}^q \sum_{b=1}^q J_{ij}(a, b) \quad (3.5)$$

Then each matrix $J(a, b)$ is symmetrized by simply averaging it with its transpose:

$$\hat{J}(a, b) \leftarrow \frac{1}{2} (\hat{J}(a, b) + \hat{J}(a, b)^T) \quad \forall a, b \quad (3.6)$$

Finally, a contact map is obtained by normalizing \hat{J} over the states and applying an average product correction w.r.t. the sites. plmDCA shows remarkable performance on diverse sets of proteins with running times competitive with mean-field DCA.

3.1.4 Gaussian Direct Coupling Analysis

plmDCA is of state-of-the-art performance, but still requires high computational resources. An alternative method is to use GaussDCA [7], which does exact inference without having recourse to iterative algorithms.

In GaussDCA, each variable $x_i \in \{0, 1\}$ indicates whether residue located at locus $i \% L \in \{1, \dots, L\}$ is of amino acid type $i/L \in \{1, \dots, 22\}$. With such a formalism, each protein is described as a vector $x \in \{0, 1\}^{22L}$. The key assumption at the core

of the method is to approximate each binary variable x_i by a continuous Gaussian variable. Let m be the number of sequences in a MSA, $X \in \{0, 1\}^{m \times 22L}$ be the matrix representation of the MSA, and μ, \bar{x} be respectively, the theoretical and empirical mean vectors associated to X . The empirical covariance matrix of X is given by:

$$\bar{C}_{ij}(X, \mu) = C_{ij}(X, \mu) = \frac{1}{m} \sum_{k=1}^m (x_i^k - \mu_i)(x_j^k - \mu_j) \quad (3.7)$$

Similarly to PSICOV [47], evolutionary couplings are detected by keeping track of direct interactions between variables of the system, which can be realized by computing the precision matrix, also known as the inverse of the covariance matrix. When matrix \bar{C} is not rank deficient, maximum log-likelihood is attained by setting the theoretical covariance matrix Σ to \bar{C} . However, the empirical covariance matrix rarely has full rank due to the limited number of effective sequences in MSAs. The suggested solution is to provide a prior distribution on positive-definite matrices and perform exact Bayesian inference to find an invertible covariance matrix.

3.1.4.1 Bayesian inference

In Bayesian inference, a hypothesis is favoured over others based on its posterior probability, which is proportional to the product of the data log-likelihood under this hypothesis and its prior probability. This relation holds in the parametric formulation of Bayes' rule:

$$P(\theta|X, \alpha) = \frac{P(X, \theta|\alpha)P(\theta|\alpha)}{P(X|\alpha)} \propto P(X, \theta|\alpha)P(\theta|\alpha) \quad (3.8)$$

$P(\theta|\alpha)$ is a prior distribution, hence a distribution over the parameter space without any knowledge about the data X . As more and more data becomes available, the newly observed samples can be incorporated to the model through the likelihood $P(X, \theta|\alpha)$. The likelihood measures how strongly the data is explained by the model, given a set of parameter values. $P(X|\alpha)$ is called evidence or marginal likelihood because it is equal to the data likelihood marginalized over the parameters θ . $P(\theta|X, \alpha)$ is the posterior distribution, a probability of some parameters given the observed data.

The marginal likelihood is a constant term that varies only with the hyper-parameters α . Therefore, the maximum a posteriori estimation (MAP) is simply the maximum product between the likelihood and the prior.

3.1.4.2 Prior distribution

A reasonable choice for the prior distribution over μ and Σ is the normal-inverse-Wishart (NIW) distribution, which is known to be conjugate prior for the Gaussian log-likelihood. The prior and the posterior are said to be conjugate if they are of the same form. In the case of NIW prior, the posterior is also a NIW distribution. The prior is defined as the product $P(\mu, \Sigma) = P(\mu|\Sigma) P(\Sigma)$, where the prior of the mean vector is defined as a multivariate Gaussian distribution:

$$P(\mu|\Sigma) = \left(\frac{2\pi}{\kappa}\right)^{-\frac{m}{2}} |\Sigma|^{-\frac{1}{2}} \exp\left(\frac{\kappa}{2}(\mu - \eta)^T \Sigma^{-1}(\mu - \eta)\right) \quad (3.9)$$

where κ is the number of prior measurements and η the prior mean.

Let's note that the Gaussian distribution is conjugate to itself. The prior over positive-definite matrices (we are interested in invertible covariance matrices exclusively) is defined by the NIW distribution:

$$P(\Sigma) = \frac{1}{Z} |\Sigma|^{-\frac{\nu+m+1}{2}} \exp\left(-\frac{1}{2} \text{Tr } \Lambda \Sigma^{-1}\right) \quad (3.10)$$

where $Z = 2^{\frac{\nu m}{2}} \pi^{\frac{m(m-1)}{4}} |\Lambda|^{-\frac{\nu}{2}} \prod_{k=1}^m \Gamma\left(\frac{\nu+1-k}{2}\right)$

where Γ is Euler's Gamma function, ν is the degree of freedom and Λ is a parameter matrix called scale matrix.

3.1.4.3 Computing the MAP covariance matrix

The mean values for the NIW prior are known and equal to η and $\frac{1}{\nu-m-1} \Lambda$, respectively. Assuming η' , ν' and Λ' are the corresponding parameters in the NIW posterior, the mean values are given by η' and $(1/(\nu' - m - 1))\Lambda'$, respectively. In particular, the matrix Λ' that allows the posterior to be conjugate with the prior is given by the following relation:

$$\Lambda' = \Lambda + n\bar{C} + \frac{\kappa m}{\kappa + m}(\bar{x} - \eta)(\bar{x} - \eta)^T \quad (3.11)$$

Finally, the average covariance matrix is computed as:

$$\begin{aligned} \Sigma &= \frac{\Lambda'}{\nu' - m - 1} \\ &= \frac{\Lambda + n\bar{C} + \frac{\kappa n}{\kappa + n}(\hat{x} - \eta)^T(\bar{x} - \eta)}{\nu + n - m - 1} \\ &= \lambda\Lambda + (1 - \lambda)\bar{C} + \lambda(1 - \lambda)(\bar{x} - \eta)^T(\bar{x} - \eta) \end{aligned}$$

As a viable choice for the hyper-parameter matrix Λ , one could choose the most trivial one to avoid arbitrary choices. To this aim, the covariance matrix corresponding to a uniform multivariate distribution is used.

3.2 PSICOV

3.2.1 Gaussian graphical models

PSICOV relies on a Gaussian graphical modelization. Such models are represented by undirected graphs that satisfy the pairwise Markov property: two nodes are not connected by an edge if they are conditionally independent. Let $G = (V, E)$ be a graph and let each node k be associated to a random variable X_k . Conditional independence between variables X_u and X_v is formally defined as follows:

$$X_u \perp\!\!\!\perp X_v | X_{V \setminus \{u,v\}} \text{ if } \{u, v\} \notin E \quad (3.12)$$

If the graph is not complete, then there exists pairs of conditionally independent variables in the model, resulting in sparsity patterns in the underlying precision matrix. When the covariance matrix of the multivariate Gaussian model is rank deficient, a solution for the precision matrix has to be found with a high prior towards sparse matrices. Such a solution can be achieved through regularization: the solution that follows in the next section relies on pseudo-inverse computation with L_1 regularization.

3.2.2 Evolutionary couplings

The motivation behind PSICOV [47] is that residue contacts produce constraints on the types of residues in the protein at certain pairs of sites: two residues involved in a contact should have complementary physicochemical properties. To capture the covariation of residue types for any pair of sites, random variable $X_{ia} : \Omega \rightarrow \{0, 1\}$ is defined, where Ω is a set of two possible outcomes (either the amino acid at site i is of type a or not). In particular, value $x_{ia}^{(k)}$ is equal to one if the k th sequence in the given MSA has a residue of type a at site i . With this knowledge, the sample covariance matrix $S \in \mathbb{R}^{21L \times 21L}$ can then be computed by:

$$S_{ij}^{ab} = \frac{1}{n} \sum_{k=1}^n (x_{ia}^{(k)} - \bar{x}_{ia})(x_{jb}^{(k)} - \bar{x}_{bj}) \quad (3.13)$$

where n is the number of homologous sequences, S_{ij} is a covariance matrix and S_{ij}^{ab} is the sequence identity covariance from amino acid at site i (being in state a) to amino acid at site j (being in state b). It is noteworthy that such a covariance matrix S_{ij} is neither positive semidefinite nor symmetric and is, thus, substantially different from regular sample covariance matrices. Precision matrix Θ_{ij} is defined as the inverse of covariance matrix S_{ij} , but there is no guarantee that such an inverse can be computed directly. Indeed, each submatrix S_{ij} is very likely to be singular due to the absence of correlations for some residue types that are rarely (or never) observed at given sites.




Figure 3.1: Illustration of a small system with only 3 sites and 2 amino acid types A and B: (a) Gaussian graphical model representing the system. (b) Sparsity pattern of the estimated inverse covariance matrix. Cells associated to zero variables are drawn in black. (c) Corresponding MSA.

3.2.3 Inference

The solution chosen here is to estimate sparse inverse covariance matrices instead, by having recourse to Graphical Lasso algorithm [32]. The method assumes that observations follow a multivariate normal distribution characterized by the following probability density function:

$$f(x) = \frac{1}{\sqrt{(2\pi)^d |\Sigma|}} \exp\left(-\frac{1}{2}(x - \mu)^T \Sigma^{-1}(x - \mu)\right) \quad (3.14)$$

where $d = 21L$ is the number of components in vector x , μ is the theoretical mean and Σ the theoretical covariance matrix. Let's express the log-likelihood of the data as a function of the inverse theoretical matrix Θ :

$$\begin{aligned} \log L(\Theta) &= \sum_{k=1}^n \log f(x^{(k)}) \\ &= \sum_{k=1}^n \left(-\frac{1}{2}(x^{(k)} - \mu)^T \Sigma^{-1}(x^{(k)} - \mu) - \log \sqrt{(2\pi)^d |\Sigma|} \right) \\ &= \sum_{k=1}^n \left(-\frac{1}{2}(x^{(k)} - \mu)^T \Theta (x^{(k)} - \mu) - \log \sqrt{(2\pi)^d \frac{1}{|\Theta|}} \right) \\ &= -\frac{1}{2} \sum_{k=1}^n (x^{(k)} - \mu)^T \Theta (x^{(k)} - \mu) - \frac{1}{2} \sum_{k=1}^n \log \left((2\pi)^d \frac{1}{|\Theta|} \right) \\ &= -\frac{n}{2} \text{Tr}(S\Theta) + \frac{n}{2} \log |\Theta| - \frac{1}{2} \sum_{k=1}^n \log (2\pi)^d \end{aligned} \quad (3.15)$$

The latter expression holds because the empirical mean \bar{x} is equal to μ for any Σ . The

objective function of Graphical Lasso is simply the log-likelihood penalized with L_1 norm. L_1 regularization is used instead of L_2 because of its non-asymptotic behaviour and therefore its ability to produce more zeroes among the parameter values. The solution to the optimization problem is described as:

$$\begin{aligned}\hat{\Theta} &= \underset{\Theta}{\operatorname{argmax}} \quad \log L(\Theta) - \rho' \|\Theta\|_1 \\ &= \underset{\Theta}{\operatorname{argmax}} \quad -\frac{n}{2} \operatorname{Tr}(S\Theta) + \frac{n}{2} \log |\Theta| - \frac{1}{2} \sum_{k=1}^n \log(2\pi)^d - \rho' \|\Theta\|_1 \\ &= \underset{\Theta}{\operatorname{argmax}} \quad -\operatorname{Tr}(S\Theta) + \log |\Theta| - \rho \|\Theta\|_1\end{aligned}\tag{3.16}$$

The L_1 norm $\|\Theta\|_1$ is the sum of absolute values of the elements in Θ . ρ is the regularization parameter and ρ' is syntactic sugar for denoting the same parameter before division.

In the second version of PSICOV, the objective function present in equation 3.16 is being optimized with the GLassoFast algorithm [91], making the predictor more efficient than in its first version.

The predicted contact map is finally obtained after applying the two following processing steps:

- The score $S_{ij}^{contact}$ between residues i and j is calculated as the L_1 norm of sub-matrix Θ_{ij} :
- The average product correction introduced by Dunn et al. [22] is applied to the resulting scores in order to better approximate the background mutual information between sites i and j :

$$PC_{ij} = S_{ij}^{contact} - \frac{\bar{S}_{i-}^{contact} \bar{S}_{-j}^{contact}}{\bar{S}^{contact}}\tag{3.17}$$

Based on the formalism of the PSICOV study, $\bar{S}_{i-}^{contact}$ is the norm of row i of $S^{contact}$ for all columns except i , and $\bar{S}_{-j}^{contact}$ is the norm of column j for all rows except j .

According to Sun et al. [90], the solutions found by PSICOV may not be suitable for some proteins. Furthermore, the best solution could be far from the global optimum found by Graphical Lasso since the search space is huge. The suggested solution is to predict multiple contact maps and promote diversity among them by adding a new penalty term.

$$\begin{aligned}\min_{\Theta^{(m+1)}} \quad & \frac{1}{2} \operatorname{Tr}(S\Theta) + \frac{1}{2} \sum_{k=1}^n \log(2\pi)^d - \frac{1}{2} \log |\Theta| \\ \text{s.t.} \quad & d(\Theta^{(k)}, \Theta^{(m+1)}) \geq \epsilon, \quad k = 1, \dots, m\end{aligned}\tag{3.18}$$

In that last equation, the distance constraint function d is a convex and differentiable function defined as:

$$d(\Theta^{(k)}, \Theta) = - \sum_{i,j} \delta_0(\theta^{(k)})_{i,j} |\Theta_{i,j}| \quad (3.19)$$

where $\delta_0 : \mathbb{R}^{21L \times 21L} \rightarrow \mathbb{R}^{L \times L}$ is such that $\delta_0(\Theta_{i,j})$ takes value 1 if submatrix $\Theta_{i,j}$ is 0, and -1 otherwise. The objective function is optimized using a second-order method and updating the solution in the Newton direction at each iteration. Finally, a contact map is obtained by selecting the submatrices among the solutions according to their sparsity. More specifically, submatrices are selected according to their nuclear norm, hence the sum of their singular values. The authors claim that nuclear norm is better than L_1 norm at taking into account the sparsity patterns of the different submatrices.

3.3 Deep learning in Protein Contact Prediction

3.3.1 Input features

3.3.1.1 Global features

3.3.1.1.1 Protein length Protein length is computed as the number of residues in the protein sequence. It provides complementary information that convolutional layers cannot capture. Indeed, fully-connected neural networks have the ability to handle arbitrary-sized features maps at the cost of not knowing the dimensionality of their inputs. Injecting protein length as a supplementary global feature may help the model to infer the maximum distance in long-range contacts.

3.3.1.1.2 Effective number of sequences The number of effective sequences is equal, w.r.t. to a given threshold, to the number of non-redundant sequences in the set of homologous sequences. It provides a bound on the potential performance of the DCA methods involved in the pipeline.

$$M_{eff} = \sum_{i=1}^m w_i = \sum_{i=1}^m \frac{1}{|\{r(s_i, s_j) \geq \tau \quad \forall j \in \{1, \dots, m\}\}|} \quad (3.20)$$

w_i is the weight associated to homologous sequence i . $r(s_i, s_j)$ is the identity rate between sequences s_i and s_j , in other words the number of matching residues (including gaps) divided by the total number of positions, which is assumed to be equal in both sequences.

3.3.1.2 1-dimensional features

3.3.1.2.1 One-hot encoded sequence Given an amino acid sequence $\{s_1, s_2, \dots, s_L\}$ of size L , the one-hot encoded sequence is a matrix $X \in \{0, 1\}^{L \times 21}$ where x_{ia} is one if $s_i = a$ and zero otherwise.

3.3.1.2.2 Solvent accessibility prediction Solvent accessibility or solvent-accessible surface is the surface of a molecule that is accessible to a solvent. Such feature contains indirect information about the protein structure since residues with low accessibility are more likely to be internal to the protein and vice versa. When atomic coordinates are available, solvent accessibility can be computed exactly using the rolling ball algorithm from Shrake and Rupley [85], as illustrated by figure 3.2. Since protein structure is not available for target proteins of the test set, supervised models rely on predicted solvent accessibility instead.




Figure 3.2: Accessible surface, obtained by "rolling" a probe sphere (a molecule of solvent, colored in orange) on the Van der Waals surface of a biomolecule (colored in blue).

The RaptorX-Property server [97] is publicly available and provides predictions for relative solvent accessibility (RSA). The latter is defined as the predicted accessible surface of a residue divided by the maximum possible solvent accessibility for that amino acid, which is more convenient more machine learning since it does not a normalization step. RSA prediction is 3-state. A residue can be either:

- Buried (B), when RSA is below 10%.
- Intermediate (I), when RSA is between 10% and 40%.
- Exposed (E), when RSA is above 40%.

3.3.1.2.3 Predicted secondary structure prediction As defined by DSSP, there are 8 different states for encoding secondary structure: H (alpha-helix), G (310 helix), I (pi-helix), E (beta-strand), B (beta-bridge), T (beta-turn), S (high curvature loop) and L (irregular loop). In practice, secondary structure prediction is either 3-state or 8-state [97]. The eight states are: Alpha-helix, 310 helix, Pi-helix, Beta-strand, Beta-bridge, Beta-turn, High curvature loop, Irregular loop.

When the number of labels is limited to three, the predictor only focuses on beta-strands (E), alpha-helices (H) and a third state which is the union of the six remaining states (C). Formally, a m -state secondary structure prediction is a matrix $S \in [0, 1]^{L \times m}$ where element $S_{i,j}$ is the probability of residue i being in conformation j , with each row summing to one.

3.3.1.2.4 Region disorder prediction Region disorder prediction is a vector $D \in [0, 1]^L$ where D_i is the probability of residue i being in a region of missing residues in the X-ray 3D structure. Residues with high probability are said to be disordered. Such features are also made publicly available by the RaptorX-Property server [97].

3.3.1.2.5 Amino acid frequencies Amino acid frequencies are position-specific features that can be efficiently computed. Let $S \in \{0, \dots, 22\}^{M \times L}$ be a Multiple se-

quence alignment matrix containing M sequences aligned to a target sequence of length L . Then amino acid frequencies can be arranged in a matrix $F \in \mathbb{R}^{L \times 22}$ where element F_{ia} is computed as follows:

$$F_{ia} = \frac{1}{M} \sum_{k=1}^M \delta(S_{ki}, a) \quad (3.21)$$

3.3.1.2.6 Position-Specific Scoring Matrix The position-specific scoring matrix [46] is computed on the whole multiple sequence alignment as well. It is of shape $F \in \mathbb{R}^{L \times 22}$ like the amino acid frequency matrix but is more informative as it summarizes common patterns present in the alignment.

3.3.1.2.7 Atchley factors Atchley factors are residue-specific constant values resulting from the factor analysis of Atchley [6]. There are 5 factors:

- Factor I is a polarity index but incorporates many variables like polarity, hydrophobicity, the covariance of the proportion of exposed residues (with low surface accessibility), the number of hydrogen bond donors, etc.
- Factor II is a secondary structure factor that is inversely related to the propensity of the amino acid to be in a specific structural state.
- Factor III is related to the residue size: residue volume, average volume when buried, residue weight, etc.
- Factor IV is related to the amino acid composition: number of codons, etc.
- Factor V is an electrostatic charge index.

3.3.1.2.8 Self-information In information theory, self-information is the amount of information, in bits, obtained by observing a random variable. In particular, let $x_{ij} \in \{0, 1\}$ be a binary variable indicating the presence of an amino acid of type j at site i . The self-information suggested by Michel et al [64] can be formalized with the following equation:

$$I_{ij} = \log_2(p_{ij}/\langle p_i \rangle) \quad (3.22)$$

where p_{ij} is the probability of observing amino acid j at site i among all residues of given MSA, and $\langle p_j \rangle$ is the frequency of amino acid j in the Uniref50 dataset.

3.3.1.2.9 Partial entropies Partial entropies are computed in [64] with the following formula:

$$S_i = p_i \log_2(p_i/\langle p_i \rangle) \quad (3.23)$$

where p_i is the probability of finding a particular amino acid or a gap.

3.3.1.3 2-dimensional features

3.3.1.3.1 Mutual Information and Normalized Mutual Information Following the formalism described in [64], MI is described as:

$$MI(x, y) = \sum_{x, y} P(x, y) \log \left(\frac{P(x, y)}{P(x) \cdot P(y)} \right) \quad (3.24)$$

where $P(x)$ is the probability associated to variable x . Normalized mutual information (NMI) is then defined as:

$$NMI(x, y) = \frac{MI(x, y)}{\sqrt{S(x) \cdot S(y)}} \quad (3.25)$$

where $S(x)$ is the entropy of variable x . Average product correction is applied to both MI and NMI.

3.3.1.3.2 Cross-entropy Cross-entropy is computed in [64] using the following formula:

$$H(x, y) = S(x) + S(y) - MI(x, y) \quad (3.26)$$

3.3.1.3.3 Evolutionary couplings Predictions from GaussDCA, plmDCA or PSICOV are the most discriminative features used in supervised methods. Only top predicted contacts are informative, but deep architectures help refining predicted couplings into high-quality contact maps.

3.3.1.3.4 Covariance Covariance matrices are computed as in equation 3.13 from the section about Gaussian graphical models (see section 3.2.2). In PSICOV [47], the inferred covariance matrix is averaged across the dimension of amino acid types. In DeepCov [48], the full matrix is used as input to the supervised model.

3.3.2 Features for the proposed approach

The lines which will follow are a discussion about the features to use as inputs to the model. At the outset, it should be noted that all models rely on two-dimensional features. Those features are either predictions from another predictor, covariance matrices or correlated mutations. However, using intermediary predictions requires installing additional software. PSICOV is written in C and can be recompiled. The official implementation of plmDCA was originally made in Matlab, forcing researchers to add a whole raft of external tools. plmDCA, as well as GaussDCA now have a Julia implementation, making them more accessible. The proposed approach incorporates all three predictors in its pipeline.

Features		DeepCov					DeepContact	PConsC4	DNCON2	RaptorX	Proposed method
Global	Protein length								×		×
	Meff								×		×
1-dimensional	Column log-entropy						×				
	Predictors stdv						×				
	Encoded sequence							×			×
	Solvent accessibility								×	×	×
	Predicted SS								×	×	×
	AA frequencies						×				
	PSSM								×	×	
	Atchley factors										
	Self-information							×			×
	Partial entropies							×			×
2-dimensional	Mutual Information (MI)							×		×	×
	Normalized MI						×	×			×
	Cross-entropy						×	×			×
	Contact potential									×	
	EVFold						×				
	CCMPred						×			×	
	plmDCA								×		
	GaussDCA							×			×
	PSICOV										
	Covariance										

Table 3.1: Features used in state-of-the-art deep learning approaches. Feature extraction methods that rely on external tools (excluding MSA tools) are highlighted in bold.

3.3.3 Recurrent networks

Recurrent neural networks [16], and more specifically their extension called long short term-memories (LSTMs) [41] have been considered for designing contact predictors due to their ability to accumulate long-range information along proteins. LSTM training is very different from the standard backpropagation algorithm introduced in section 2.2.6 due to the presence of feedback connections. SPOT-Contact [38] uses both residual convolutional networks and LSTMs to predict contact maps, and has shown to outperform many state-of-the-art models on CASP12 targets [71].

3.3.4 Fully-convolutional networks

As a reminder, fully-convolutional networks are neural networks capable of handling variable sized input. Therefore, the dynamic input dimensions cannot be processed by a fully-connected neural layer. An example of such an architecture is DeepCov [48], a deep neural network composed only of 2D convolutional layers and an additional maxout input layer. A maxout layer is made of a convolutional layer, followed by a max-pooling operation. Dimensions of intermediate feature maps are preserved by convolutions using a stride of one and a "same-padding". In most deep learning libraries, padding of type "same" maintains the sizes of spatial (convoluted) dimensions. In DeepCov, the only features are the covariance matrices as computed in equation 3.13. Couplings matrices predicted by DCA methods can be fed as input to the model instead of covariance matrices, as shown in plmConv study [34].

The approach described in the DNCON2 paper [48] also implements convolutional layers with dynamic spatial dimensions. Additionally to that, the fuzziness of residue contact definition is handled by training one fully-convolutional neural network per contact threshold. More specifically, five networks are trained to output contact maps at 6, 7.5, 8, 8.5 and 10 Å thresholds, respectively. These five are stacked on top of a sixth network in charge of refining and combining the predictions into a final contact map at a threshold of 8 Å.

3.3.5 Residual Networks (ResNets)

As described in section 2.2.6 about backpropagation, the loss gradient with respect to the parameters of a specific layer is computed as the product of many other mathematical entities (vectors, scalars, matrices, etc.), and the number of factors in such a product grows linearly with the number of operations applied after current layer. When this number is too large, some layers may be updated with numerically unstable gradients.

A widely used solution is to add residual connections [40] to the architecture. The latter can thus no longer be viewed as a regular composition of functions and must take into account the residual mapping at the end of each residual block. The output $Y^{(r)}$ of residual block k should now be formalized with a more general form:

$$Y^{(r)} = f(X^{(r)}, \{W^{(p)}\}_p) \quad (3.27)$$

where $W^{(p)}$ are the weights of a layer p in block k . Residual mappings can be imple-

mented in several ways and figure 3.3 illustrates one of them.



Figure 3.3: Illustration of a residual connection in a convolution neural network. For the element-wise sum to work, the input and output of the residual block are required to be of the same shape.

3.3.6 Deep fully-convolutional residual networks

Most successful methods rely on very deep architectures with residual mappings. One-dimensional features are processed by one-dimensional residual network before being concatenated with two-dimensional features. The resulting tensor is the input of a two-dimensional residual network. Examples of such approach is DeepContact [60] and the state-of-the-art RaptorX-Contact predictor [99]. Both methods rely on CCMPred contact prediction, solvent accessibility and secondary structure prediction. Additionally to CCMPred, DeepContact incorporates EVFold predictions together with the rest of the two-dimensional features. Also, global features (e.g. the number of effective sequences) are tiled and concatenated with other features: this does not impact the fully-convolutional property of DeepContact because global features are invariant to the protein length. The largest difference in the two methods lies in the depth of the networks: 9 convolutional layers for DeepContact and 60-70 layers for RaptorX-contact. It must be noted that RaptorX-contact architecture is not fully-convolutional since zero-padding is applied to feature maps when more than one protein is processed in a batch.

3.3.7 U-net architecture

In PconsC4 [64], the model is partly built on top of the U-net architecture [79], and trained on a set of 2891 proteins retrieved from PDB. Features are divided in one-dimensional inputs including one-hot-encoded amino acids, self-information and partial entropies, and two-dimensional inputs including mutual information, normalized mutual information and cross-entropy. For both mutual information and normalized mutual information, average product correction is applied. One-dimensional features are convoluted through one-dimensional residual networks, concatenated and finally reshaped to two-dimensional maps with an outer product. After reshaping, the intermediate feature maps are concatenated with the two-dimensional features and the whole is fed as input to the U-net architecture.

Models designed for semantic segmentation problems (including PCP) have been shown to gain significant performance when additional connections are allowed between layers close to the network’s input and layers close to the network’s output [43]. U-net architectures develop this idea: they have a somehow symmetric structure made of two

sub-networks that compress and decompress the information, respectively. The first sub-network successively alternates between convolutional transformations and max-pooling, while the second sub-network alternates between convolutional transformations and transposed convolutions (also called upsampling). Similarly to autoencoders, the features maps processed the middle layers or of much smaller dimensionality than input or output tensors. To prevent the whole architecture from forgetting contextual information due to compression, shortcut connections are added between tensors of identical dimensionality. Contrary to residual networks, these shortcut connections are implemented by concatenation instead of addition. Due to the max-pooling and upsampling operations, the width and height of input images should preferably be a power of two. Because PconsC4 model outputs entire contact maps, and because proteins are of variable length by nature, input feature maps of a particular protein are zero-padded to the smallest power of two that is larger than the protein length. Finally, PconsC4 architecture does not only predict a contact map but predicts multiple contact maps at different distance thresholds.

3.3.8 Dense networks (DenseNets)

Dense networks extend the idea of residual networks by allowing residual mappings between all layers, resulting in an ergodic topology. In DenseNets [43], shortcut connections are implemented with a concatenation operation instead of addition. Nevertheless, spatial dimensions are still constrained to be identical between the source and the destination of a shortcut connection. Figure 3.4 illustrates a dense block composed of 2 convolution layers: it is shown how connections are allowed between all entry points and layer outputs.




Figure 3.4: Illustration of dense blocks from the DenseNet paper [43]. All layer entry points and outputs belonging to a dense block are connected by residual mappings.

Each unit of a dense block is the composition of batch normalization, activation and convolution operations.

3.3.9 TiramProt

TiramProt is based on the Tiramisu architecture [93] which tries to combine the ideas of both U-net and DenseNet. Like U-net architectures, Tiramisu applies max-pooling for reducing dimensionality and upsampling to restore the input dimensionality. Each dense block output from the downsampling part of the network is, by symmetry, connected to the entry point of the corresponding dense block in the upsampling part of the network. As in U-net, these shortcut connections are implemented by feature map concatenation. Training set and features are the same as in the Pconsc4 study.

3.3.10 DeepConPred2

DeepConPred2 [20] is based on a problem-specific architecture made of three modules. First module consists of a Deep Belief Network (DBN) that predicts contacts between secondary structures from CCMPred coevolutionary information. A DBN is a graphical model implemented as a stack of unsupervised building blocks such as autoencoders or restricted Boltzmann machines. The output of the first module, together with solvent accessibility and secondary structure prediction, is fed as input to each of the DBNs composing the second module: one for short-range contacts, one for medium-range contacts and three for long-range contacts (taken from previous study [103]). Each of these components are used to predict actual residue contacts. Then, each DBN output is fed as input to one of the ResNets of the third module. Each ResNet has 50-80 convolutional layers with Leaky-ReLU activation functions.

3.3.11 Properties of DL approaches

PCP methods that have just been described are summarized in table 3.2 according to some indicators such as training set size, fully-convolutional and fuzzy aspects of the approach and depth of the architecture. RaptorX-contact has not been counted as a fully-convolutional approach since zero-padding is used for batches of more than one protein. Depth of the network has been measured as the number of non-linearities, that is to say the number of layers or blocks preceding an activation function. Additionally, the table shows which methods incorporate contacts defined at different distance thresholds.

	Training set size	Fully-convolutional	# non-linearities	Fuzzy
DNCON2	1230	-	6	⊤
DeepCov	6003	⊤	14	⊥
plmConv	231	⊥	4	⊥
DeepConPred2	3443	-	50-80	⊥
DeepContact	-	⊤	9	⊥
PconsC4	2891	⊥	-	⊤
TiramiProt	2891	⊥	16	⊤
RaptorX-Contact	~ 6000	⊥	60-70	⊥

Table 3.2: Overview of different deep learning models for PCP: number of proteins in the training set, whether the architecture is fully-convolutional, how deep (measured in activation functions) the model is, and whether it learns from contact maps defined with multiple distance thresholds.

In PconsC4 and TiramiProt, a single network predicts multiple contact maps at the same time, while DNCON2 uses multiple networks to each predict a contact map at a different threshold.

3.4 Contact-assisted protein folding

Static protein structure can be recovered with a high resolution solely based on a few true residue contacts [50]. Because predicted contact maps are fuzzy by nature and may contain false contacts, only top $\alpha \cdot L$ contacts (residue pairs associated to highest predicted probabilities) are usually considered.

In the study of Chelvanayagam et al. [14], protein structure is modelled in a distance geometry setting using Gaussian restraints with empirically known mean and variance. Let $E(x_i, x_j)$ be the Euclidean distance between residues i and j , where residues are being represented by the center of their respective C_β (or C_α) atoms. Let $D_{i,j}$ and $V_{i,j}$ be the average Euclidean distance and variance associated with the Gaussian constraint between residues i and j . Under the assumption of independence between residue pairs, maximizing the log-likelihood reduces to minimizing the following objective function:

$$\sum_{i=1}^{n-1} \sum_{j=i+1}^n \frac{(E(x_i, x_j) - D_{i,j})^2}{V_{i,j}} \quad (3.28)$$

where n is the protein size in residues. Constraints are chosen according to predicted secondary structure and surface accessibility. By default, all residues are assumed to be separated by an average distance of 120 Å, with a high variance (120 Å²) to allow flexibility. For each residue pair, these prior parameters are replaced by a more accurate constraint when sufficient information is available. For example, residues with low surface accessibility are assumed to be separated by a distance of 7.5 Å from the center of mass of the protein, while residues with high surface accessibility or assigned a higher average distance of 12 Å. The Euclidean coordinates of the center of mass are additional variables to be added to the model. Residues participating in the active site are considered to be near in space. Adjacent and almost-adjacent residues are assigned low distances with very small variance since the distance between adjacent C_α atoms is fixed but with a small variability induced by torsion angles. Helices and strands, when predicted secondary structure is available, are modelled as well. All residue pairs with a sequence separation $\in [2, 5]$ in alpha helices are being constrained, and only residue pairs with a sequence separation $\in [2, 3]$ in 3/10 helices are being constrained. Similarly, sequence separation in strands must be in the range [2, 4]. Disulfide bonds are modelled by an average distance of 5.5 Å and a variance of 0.2 Å² between cysteine residues. In sheets, center residues of adjacent strands are assigned an average distance of 4.54 Å and a variance of 0.1 Å². However, sheet topology and disulfide bonds are not always available and the solution suggested in the paper is to have recourse to sheet and disulfide combinatorics.

GDFuzz3D [76] extends these ideas without having recourse to combinatorics but introducing the idea of graph distance (length of the shortest path between two residues). It is two-stage: first residue positions are optimized by minimizing the objective function previously introduced, and then atom coordinates are being refined with MODELLER [28] in order to obtain an all-atom model. Current state-of-the-art contact-assisted methods are CONFOLD [1, 2] and the very similar predictor CoinFold [98].

Chapter 4

Materials and methods

4.1 Datasets

Seven datasets have been used in the framework of this thesis: three for training the model, one for optimizing the hyper-parameters and validating the model, and three for benchmarking. Training set is composed of the PSICOV150 dataset [47], PConsC4's training set and PConsC3's benchmark set. Validation set is composed of the 30 additional training proteins introduced by PConsC3 [88]. Three test sets have also been considered in order to make a direct comparison with the state-of-the-art RaptorX-Contact predictor. The first test set embodies the 105 protein domains from the CASP11 experiment, the second test set 76 proteins from the CAMEO project, and the third one 398 membrane proteins.

Despite the fact that these benchmark sets are relatively old, they have been privileged due to their direct availability, including the availability of the multiple sequence alignments. Indeed, direct comparison with state-of-the-art methods requires the use of the same alignments since the quality of the prediction is heavily affected by the number of homologous sequences and the tool used for aligning them.

Protein duplicates between training sets have been removed. Also, because PConsC4's training set has been created after the CASP11 targets have been released, and mostly because the study has been conducted after the publication of the RaptorX-Contact paper, an homology reduction step had to be introduced, as explained in section 4.1.6.

4.1.1 PSICOV Dataset

The PSICOV [47] dataset is composed of 150 families and associated multiple sequence alignments taken from the Pfam database, each containing more than 1000 homologous sequences and a target sequence with high-resolution (≤ 1.9 Å) X-ray crystallographic structure. Each target sequence contains exactly one copy of the Pfam domain, has a length lower than 275 and greater than 50 residues. The number of unique sequences in each multiple sequence alignment strongly varies from one family to another. AraC-like ligand binding domain (implied in DNA-binding transcription and sequence-specific DNA binding) accounts for 511 unique sequences, compared to 74 836 sequences for the response regulator receiver domain.

4.1.2 PConsC3’s benchmark set

This dataset is a subset of PConsC2’s test set, with no homology to PConsC3’s training set. Homology is reduced by ensuring that the two sets do not contain proteins sharing the same ECOD H-group [15]. ECOD H-classes [15] are used instead of similarity rates or identity rates because they are based on a hierarchical classification and can potentially give more evidence on whether two sequences are evolutionary related.

4.1.3 PConsC4’s training set

PConsC4’s training set is composed of 2891 proteins retrieved from PDB using the protein culling server PISCES [96] with a maximum sequence identity rate 20 % between sequences, a minimum resolution of 2 Å and a maximum R-factor of 0.3. Let’s note that all proteins dating from after 2016-05-01 have been removed to prevent any risk of data contamination with PConsC4’s test set.

4.1.4 Validation set

The validation set is composed of the 30 protein families used by PConsC3 as additional training proteins, which themselves are a subset of the test set of PConsC2 [89]. This dataset has no homology (no shared ECOD H-group) with the benchmark set of PConsC3, which is used as part of the training set in the present context. Therefore, these proteins allow to assess the generalization abilities of the model during validation and reduce empirical risk on the test sets.

4.1.5 Test sets

The three test sets are the 105 CASP11 targets [69, 70], 76 CAMEO old hard targets [36] and 398 membrane proteins that have been used to assess the performance of RaptorX-Contact and PConsC3.

4.1.6 Homology reduction

PConsC3 study ensures that no homology exists between its training and test sets, and PConsC4’s training set is composed only of proteins dating from before 2016-05-01. However, this is not sufficient to ensure independence between the final training set and the three different test sets. Therefore, an additional homology reduction step has been introduced.

A straightforward and fast method for reducing homology between two set of proteins is to remove proteins that have a sequence identity rate above a given threshold. As a rule of thumb, this threshold is usually set to 40%. Identity rates were computed by running Needleman-Wunsch algorithm on each pair of proteins coming from two different datasets. Score matrix was set such that exact matches give a score of 1 and any mismatch gives a penalty of -1: this approach promotes global alignments having maximum identity rates. This method has been privileged due to its efficiency and the

large number of proteins in the training set. Identity rate between the training and benchmark sets after homology reduction are summarized in table 4.1.

Identity rate	Minimum	Average	Maximum
Training set - CASP11	5.0 %	22.7 %	39.7 %
Training set - CAMEO	3.2 %	21.5 %	35.5 %
Training set - Membrane	4.7 %	22.5 %	39.5 %

Table 4.1: Identity rates between proteins from training and benchmark sets, expressed as percentages.

Proteins with following PDB IDs have been removed from the training due to high identity with some CASP11 targets: 3U6GA, 4PKMA, 4Q34A, 4Q53A, 4QRKA, 4R03A, 4R7QA, 4RGIA, 4WJIA, 5A1QA, 5FU5A. Proteins with PDB IDs 4RUA3, 4WY9A have been rejected as they were present both in the training set and CAMEO targets. Protein 4RZ9A has been removed because of its high identity rate with 4RZAA. Also, proteins with PDB IDs 4V17 and 5AG8 have been discarded as well since they have a A chain in the training set and a B chain in the test sets (with high identity between chains). Finally, proteins with IDs 1HCZA, 1HCZA, 1JH0L, 1K6LH, 1LDIA, 1QJ8A, 1QJPA, 1QJPA, 1VCRA, 1VCRA, 1VCRA, 1VCRA, 1VCRA, 1XQFA, 1XQFA, 2CUAA, 2D2CN, 2ERVA, 2HI7B, 2NQ2A, 2NQ2A, 2NR9A, 2ONKC, 2PORA, 2Q7RA, 2VDFA, 2WJRA, 2X55A, 2ZFGA, 3ANZW, 3GP6A, 3H90A, 3LW5L, 3LW5L, 3LW5L, 3M71A, 3NYMA, 3QE7A, 3QNQA, 3ZUXA, 4E1SA, 4M5BA, 4PGRA, 4QNDA, 4RYOA, 4WDCA, 4X5MA, 4X5MA, 4XU4A, 5AZBA, 5HYAA have been removed from the training set due to high identity with some membrane proteins.

The different datasets are summarized in figure 4.1.

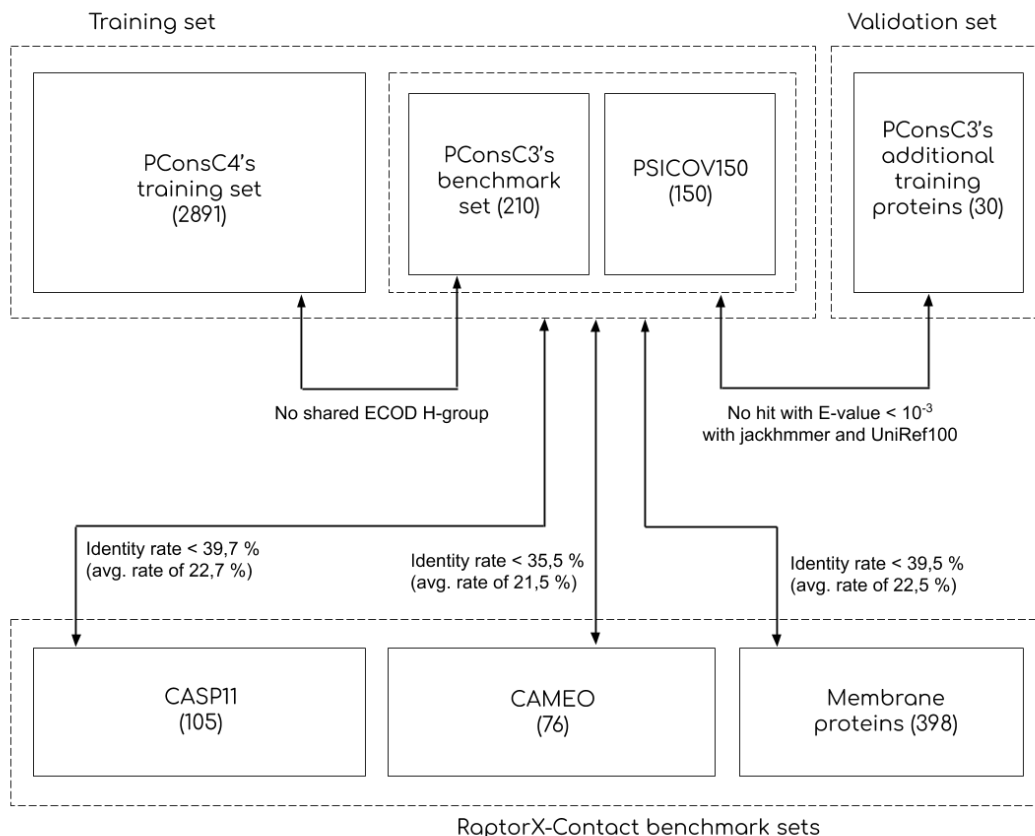


Figure 4.1: The different datasets used in this thesis, after homology reduction.

4.1.7 Feature extraction

All protein sequences, structures, multiple sequence alignments present in the three benchmark sets come from the PConsC3 server, available at the address <http://pconsc3.bioinfo.se/pred/download/> as on the date of 28th December 2018. Information available in these datasets is the following:

- Protein sequence in FASTA format
- Multiple sequence alignment obtained using HHblits on the corresponding protein family
- Native structure in PDB format
- PhyCMAP [100] intermediate predictions
- plmDCA [24] intermediate predictions
- GaussDCA [7] intermediate predictions
- Predictions made by PConsC3 [88] at each layer of the model

- CCMPred [83] predictions
- EVFold [84] predictions
- PSICOV [47] predictions
- MetaPSICOV [49] predictions

4.1.7.1 Multiple sequence alignments

Multiple sequence alignments of the PConsC4’s training set are not the same as in the original study since they have been retrieved from the RaptorX-Property server.

MSAs for PSICOV150 have been retrieved from the PConsC2 dataset. PConsC2 relied on multiple MSAs for predicting contact maps, but only the ones generated with HHblits [78] with an e-value of 10^{-4} have been considered in the present thesis.

MSAs present on the PConsC3 server have been created using HHblits (version as of the date of 26th February 2016) on the Uniprot20 database with an e-value of 1. Parameters have been set in such a way that all sequences in each of the database MSAs are aligned. These MSAs relate to the three test sets, PConsC3’s benchmark set and PConsC3’s 30 additional proteins.

4.1.7.2 Features

The Julia implementation of GaussDCA has been used with default parameters on all seven datasets. All other features, including protein length, self-information, partial entropies, mutual information, normalized mutual information, etc. are computed on-the-fly during the training process.

4.2 Summary of input features

As described in section 3.3.1, input features can be split into three categories: global, 1-dimensional and 2-dimensional features. The proposed model takes as input the protein length, the effective number of sequences in the corresponding MSA, position-specific statistics, residue pair-specific statistics, and predictions made by PSICOV, plmDCA and GaussDCA. Additionally to these features, secondary structure, solvent accessibility and region disorder are predicted by RaptorX-property server and added to the rest of 1-dimensional features. All features are being normalized before being fed as input to the model.

Category	Feature name	Dimensionality
Global	Protein length L	scalar
	Effective number of sequences M_{eff}	scalar
1-dimensional	One-hot-encoded sequence	$L \times 22$
	Self-information	$L \times 22$
	Partial entropy	$L \times 2 \cdot 22$
	Predicted secondary structure	$L \times 3$
	Solvent accessibility	$L \times 3$
	Region disorder	L
2-dimensional	Mutual information	$L \times L$
	Normalized mutual information	$L \times L$
	Cross-entropy	$L \times L$
	PSICOV predictions	$L \times L$
	GaussDCA predictions	$L \times L$
	plmDCA predictions	$L \times L$

Table 4.2: Input features of the proposed model. Global features are scalar values, whereas dimensional features are presented in the form of matrices of given shape.

4.3 Proposed architecture

Proposed predictor, as illustrated in figure 4.2, takes ideas from both PConsC4 [64] and RaptorX-Contact [99] architectures (see section 3.3). Global features, one-dimensional and two-dimensional features are being fed as input to the global, 1D and 2D modules, respectively.

- **Global module** is made of a succession of multiple fully-connected layers. Its output is repeated in order to match the dimensionality of the 1D module’s input. It must be highlighted that the fully-convolutional property of the whole network is preserved since the dimensionalities of the global module’s input and output are invariant to the protein length, despite the presence of fully-connected layers.
- **1D module** is a one-dimensional residual network. Each layer is composed of 1D convolution, 1D batch normalization and activation function. Its input is a concatenation of one-dimensional features and global module’s output. A Kronecker product is applied between the module’s output and itself in order to match the dimensionality of the 2D module’s input.
- **2D module** is a two-dimensional residual network. Each layer is composed of 2D convolution, 2D batch normalization and activation function. Its input consists of a concatenation of the 2D features and the output of the 1D module.

The output of the 2D module is the output of the whole model. Contrary to RaptorX-Contact, it does not represent a single contact map but 5 contact maps predicted at contact thresholds 6, 7.5, 8, 8.5 and 10 Å, like in PConsC4. In contrast to the other two, proposed architecture is fully-convolutional and handles variable-length proteins even with a batch size greater than one.




Figure 4.2: Proposed architecture of the deep convolutional neural network for semantic segmentation

The loss function used to train this architecture is the binary cross-entropy, a well-established metric for the training of binary classifiers. However, it has been customized in such a manner that the loss for residue contacts with a sequence separation below 6 is not backpropagated, as well as for residue pairs where the euclidean distance is not present in the native structure.

4.4 Evaluation

Protein contact maps are imbalanced by nature: they contain very few residue contacts compared to their number of residue pairs. L being the number of residues in a protein, the number of residue contacts increases linearly with L while the number of residue

pairs increases quadratically [72]. This is important because one can evaluate a model only on the L (or even less than L) predicted residue contacts the model is the most confident about. Such evaluation metric is called *best- L/k PPV* (*Positive Predictive Value*) and can be formulated as follows:

$$\text{Best-}L/k \text{ PPV} = \frac{\sum_{\substack{(i,j) \in B(L/k) \\ i-j \geq 6}} C_{i,j}}{L/k} \quad (4.1)$$

where $C_{i,j} \in \{0,1\} \forall i,j, i-j \geq 6$ are boolean values indicating a predicted contact. $B(L/k)$ is the set of L/k residue pairs with most confident (highest) predicted probabilities. Contacts under a residue distance of 6 amino acids are not considered during evaluation phase even though they are used during training phase.

Best- L/k PPV can also be split into three separated metrics: short-range, medium-range and long-range contacts. These three types of contacts can be defined with the following sequence separations:

- Short-range contacts: 6 - 12 residue separation
- Medium-range contacts: 12 - 24 residue separation
- Long-range contacts: 24+ residue separation

4.5 Hyper-parameter optimization

In order to ensure the best hyper-parameters are selected for the model that will be evaluated on the benchmark sets, the Hyperopt Python library [10] has been used to explore the hyper-parameter space and fine-tune the model on the validation set. Training and evaluating a deep neural network is very costly and, as a matter of fact, each trial point in the hyper-parameter space should be carefully selected. Techniques based on grid search do not suit the problem because they are uninformed methods.

4.5.1 Tree-structured Parzen estimators

Hyperopt provides an informed search method called Tree-structured Parzen Estimators (TPE) [11]. In Bayesian hyper-optimization, the posterior probability $P(\alpha|\mathcal{L})$ is defined as a function of the hyper-parameter vector α and the loss \mathcal{L} . Contrary to techniques based on Gaussian processes that approximates $P(\mathcal{L}|\alpha)$ directly, TPE models both posterior $P(\alpha|\mathcal{L})$ and $P(\mathcal{L})$. The prior is iteratively replaced with non-parametric densities based on generated points $\{\alpha_1, \alpha_2, \dots\}$. In this search, TPE is an informed search strategy that refines its prior as new points are observed in the hyper-parameter space. The "tree structure" is due to the way the posterior is computed.

Let f be the prediction function of the model (see section 2.2.6 about backpropagation algorithm), and let $f(\alpha) = \operatorname{argmin}_{f(\Theta, \alpha)} \ell(f(\Theta, \alpha))$ be the prediction function of a trained model that minimizes a given loss function ℓ w.r.t. a fixed hyper-parameter

vector α . Let $l(\alpha)$ be the non-parametric density function defined as a mixture of density functions centered each on an observation $\{\alpha^i\}$ for which $\mathcal{L} = c(f(\alpha^i))$ is below the threshold \mathcal{L}^* . Density function $g(\alpha)$ is defined analogously as a mixture of density functions centered each on one of the remaining observations. As described in the following equation, the density function to be used to approximate the posterior is determined according to whether the threshold \mathcal{L}^* has been exceeded.

$$P(\alpha|\mathcal{L}) = \begin{cases} l(\alpha) & \text{if } \mathcal{L} < \mathcal{L}^* \\ g(\alpha) & \text{otherwise} \end{cases} \quad (4.2)$$

The threshold \mathcal{L}^* is set as a quantile of the observed values of \mathcal{L} . The value to be optimized in TPE is the Expected Improvement (EI), measured as an integral of loss improvements weighted by the posterior itself. After applying Bayes formula to the posterior, calculus of EI becomes:

$$EI_{\mathcal{L}^*}(\alpha) = \int_{-\infty}^{\mathcal{L}^*} (\mathcal{L}^* - \mathcal{L}) P(\mathcal{L}|\alpha) d\mathcal{L} = \int_{-\infty}^{\mathcal{L}^*} (\mathcal{L}^* - \mathcal{L}) \frac{P(\alpha|\mathcal{L})P(\mathcal{L})}{P(\alpha)} d\mathcal{L} \quad (4.3)$$

In the framework of Adaptive Parzen Estimators, to each hyper-parameter is assigned a prior and a density function, and the estimator is built as a weighted mixture of them. For example, a continuous variable can be assigned:

- A uniform prior with lower bound a and upper bound b .
- A function defined as a mixture of Gaussian distributions, each centered on a point of the hyper-parameter space. The standard deviation of a particular distribution can be set as the maximum between distances to the left and right neighbors.

The density function is either $l(\alpha)$ or $g(\alpha)$ depending on whether the loss function associated to current point is below the threshold or not.

4.5.2 Search space

Table 4.3 summarizes the hyper-parameters to be optimized with Hyperopt. Each hyper-parameter follows a uniform prior distribution of categorical values.

- **Batch size:** Number of proteins to evaluate before updating the network parameters. An upper bound of 32 proteins has been chosen in order to restrain memory consumption to reasonable values. Indeed a large batch size, coupled with a high number of layers per module and large number of convolutional filters per layer can lead to huge computation graphs of size that easily exceed available memory.
- **Batch normalization:** Whether to normalize values in current batch (see section 2.2.6). Can be either true or false.
- **Track running stats:** Whether to normalize values in current batch using running means and variances. If set to false, statistics of current batch are used instead. Can be either true or false.

- **Learning rate:** Parameter that is proportional to the step length to perform in the improvement direction returned by the optimization algorithm. The domain is defined by a set of very small values to ensure local convergence. However, the smallest value does not fall below 10^{-6} to prevent the model from training for days.
- **L2 penalty:** L2 regularization parameter (see section 2.2.7.2).
- **Optimization algorithm:** Which optimization method to use for updating network parameters (see section 2.2.8).
- **Activation function:** Which non-linearity to use to produce the output of each layer (see section 2.2.5), except the final layer of the whole network where it is automatically replaced by a sigmoid function.
- **Use global module:** Whether to incorporate global features.
- **Depth:** Number of layers (either fully-connected or convolutional) in current module.
- **Filter size:** Number of adjacent residues being observed at once in a convolutional layer.
- **Number of filters:** Number of convolutional filters, hence the number of output channels in a convolutional layer.

Module	Hyper-parameter	Set of values
General	Batch size	$\{1, 2, 4, 8, 16, 32\}$
	Batch normalization	$\{\top, \perp\}$
	Track running stats	$\{\top, \perp\}$
	Learning rate	$\{10^{-3}, 10^{-4}, 10^{-5}, 10^{-6}\}$
	L2 penalty	$\{10^{-3}, 10^{-4}, 10^{-5}, 10^{-6}\}$
	Optimization algorithm	$\{\text{ADADELTA}, \text{Adagrad}, \text{Adam}\}$
	Activation function	$\{\text{ReLU}, \text{ELU}, \text{LeakyReLU}, \text{Tanh}\}$
	Use global module	$\{\top, \perp\}$
Global module	Depth	$\{2, 3, 4, 5, 6, 7, 8, 9, 10\}$
1-dimensional module	Depth	$\{2, 3, 4, 5, 6, 7, 8, 9, 10\}$
	Filter size	$\{3, 5, 7\}$
	Number of filters	$\{8, 16, 32, 64, 128\}$
2-dimensional module	Depth	$\{2, 3, 4, 5, 6, 7, 8, 9, 10\}$
	Filter size	$\{3, 5, 7\}$
	Number of filters	$\{8, 16, 32, 64, 128\}$

Table 4.3: Hyper-parameter space for the proposed architecture.

4.6 Implementation

4.6.1 Availability

Source code is available at: <https://github.com/AntoinePassemiers/Wynona>.

4.6.2 Deep learning framework

Methods and results presented in this thesis have been both implemented and produced in Python. The neural architecture has been built on top of PyTorch, which is an open source deep learning framework based on Torch [17].

PyTorch does not natively handle arbitrary-sized inputs, for example fully-convolutional neural networks cannot accept images with variable height/width. For this aim, it is necessary to add a layer of abstraction so the neural models are able to process *virtual batches* of inputs. Let's define a virtual batch as the number of samples a model has to process between each parameter update. A forward pass on a virtual batch then consists in constructing one computational graph per sample and backpropagate the gradients through each one of them separately. Once all the gradients have been computed, they are collected and averaged over the sample dimension. Pytorch allows to explicitly call the forward pass, backward pass and update procedures when needed, which eases the implementation of virtual batch processing.

Model general architecture is fully-convolutional, forcing us to use only deep learning functionalities that are invariant to individual input sizes. These are:

- Element-wise operations like activation functions: ReLU, ELU, Sigmoid, etc.
- Dropout, which preserves the dimensionality of its inputs.
- Convolution, because convolutional filters have a dimensionality that is invariant to the input size (see section 2.2.4).
- Batch normalization (see section 2.2.6).
- Many non-neural transformations like arithmetic operations, Einstein summation, Kronecker product, etc.

4.6.3 Feature extraction

Many features discussed in the present document rely on amino acid counts. Despite the fact that counting algorithms such as histograms are embarrassingly parallel and naturally suitable for multiprocessing, they cannot be efficiently vectorized. For this specific reason, the scientific computing library NumPy (which has been extensively used during the experiments) is not sufficient to extract this type of features in reasonable time. Instead, C extensions have been created using the Cython compiler [8].

Chapter 5

Results

5.1 Hyper-Parameter Optimization

In order to significantly reduce the cost of evaluation (re-training the whole model from scratch with a different set of hyper-parameters), the training set used for hyper-parameter tuning has been limited to the PSICOV150 families. Hyperopt stopped after only 32 iterations with a maximum of 48,9 % of best-L PPV on the 30 validation proteins.

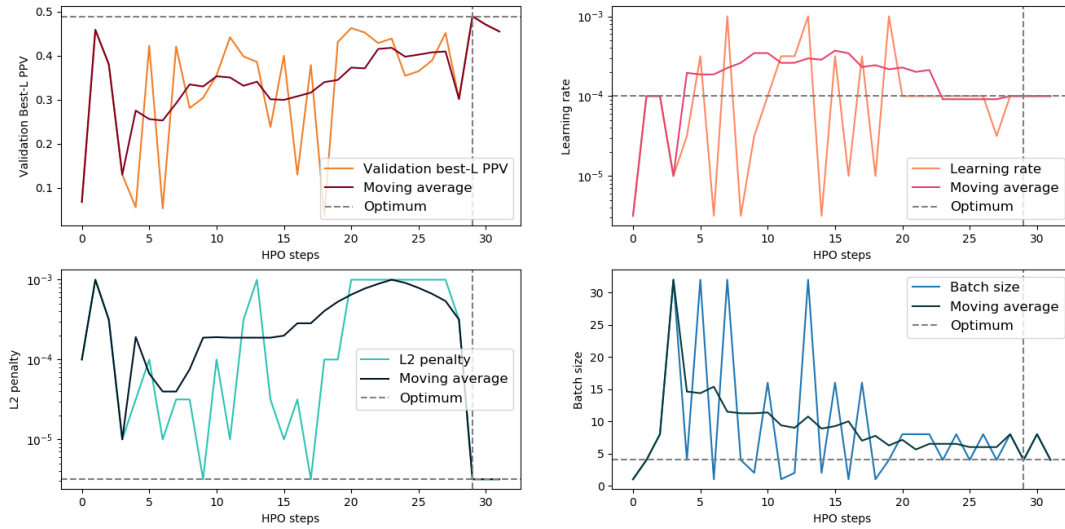


Figure 5.1: Performance and hyper-parameter values as a function of the number of Hyper-Parameter Optimization (HPO) iterations. Top left figure illustrates the optimal point whose value on x-axis is given by the HPO iteration that yields highest validation Best-L PPV. Optimal values for the learning rate, L2 penalty and batch size are denoted by dashed lines in top right, bottom left and bottom right figures, respectively.

As can be observed in top left figure in 5.1, a high-quality solution (local maximum) was found in the very first iterations, but the algorithm continued to explore and found its best solution after 30 iterations with a completely different set of parameters: L2

regularization parameter changed from 10^{-3} to $\frac{1}{2}10^{-5}$ and batch size changed from 32 proteins to only 4. Because these two solutions are very distant in the hyper-parameter space and have similar performance, it can be suggested that the latter is the search space is multimodal.

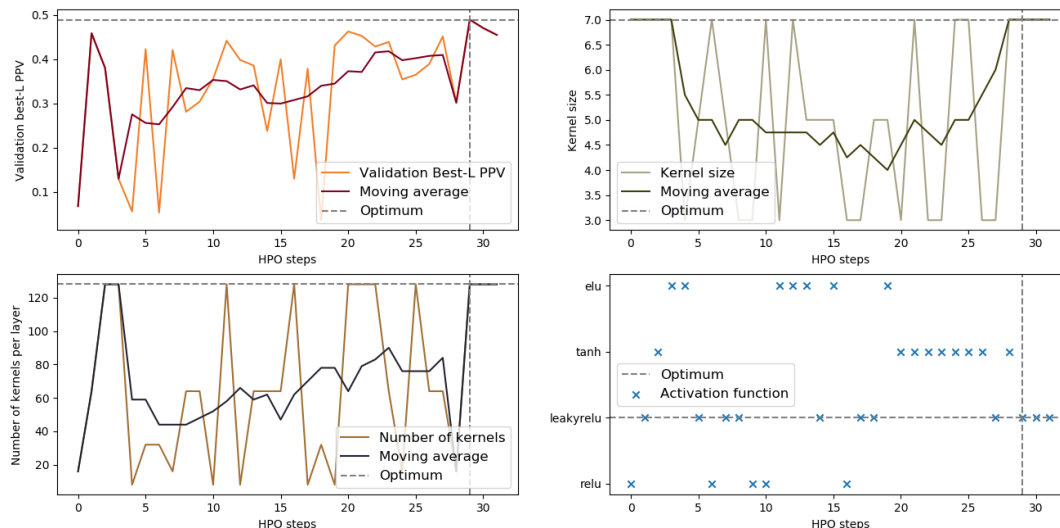


Figure 5.2: Performance and hyper-parameter values as a function of the number of Hyper-Parameter Optimization (HPO) iterations. Top left figure illustrates the optimal point whose value on x-axis is given by the HPO iteration that yields highest validation Best-L PPV. Optimal values for kernel size, number of kernels and activation function are denoted by dashed lines in top right, bottom left and bottom right figures, respectively.

Optimal hyper-parameter values are provided in table 5.1. It turned out that batch normalization is required, which is consistent with the fact that a very large network depth has been found ($3 + 18 + 18$ layers). A batch size of 4 is sufficient, which makes sense since each example is a protein of size L containing potentially $L(L-1)/2$ residue pairs and each residue pair is a statistical example *per se*.

Module	Hyper-parameter	Set of values
General	Batch size	4
	Batch normalization	\top
	Track running state	\perp
	Learning rate	10^{-4}
	L2 penalty	10^{-4}
	Parameter optimization	Adam
	Activation function	LeakyReLU
	Use global modules	\top
Global module	Depth	3
1-dimensional module	Depth	18
	Filter size	7
	Number of filters	128
2-dimensional module	Depth	18
	Filter size	7
	Number of filters	128

Table 5.1: Set of hyper-parameter values obtained at optimal point.

A L2 regularization parameter of 10^{-4} has been found. However, this value has been replaced by 10^{-5} for training the final model on the whole training set since the latter contains much more proteins than PSICOV150 alone, alleviating the need for regularization.




Figure 5.3: Binary Cross-Entropy loss and best-L PPV computed on a rolling window of batches, related to the training phase of the model with best hyper-parameters.

Convergence of the final model on PSICOV150 is shown in figure 5.3. It can be noticed that binary cross-entropy is directly and inversely related to the training PPV. The neural network takes approximately 850 batches to converge. However, when trained on the whole training set, it takes approximately 16 000 batches and 4 days before the

algorithm converges.

5.2 Model evaluation on the different benchmark sets

Performance measured in table 5.2 is the PPV obtained by only considering the top L predicted probabilities as predicted contacts (see section 4.4 for more details about the evaluation metrics). In this way, the evaluation is based only on the contacts the predictive model is the most confident about. As can be observed, best-L PPV is significantly lower for short range contacts, regardless of the method used. Indeed, the number of residue pairs having a sequence separation between 6 Å and 12 Å is much smaller than for a separation between 12 Å and 24 Å (medium range) or above 24 Å (long range). Furthermore, proteins with few or no short range native contacts will obviously yield a lot of false positives. Short range best-L PPV is thus computed on the basis of much less confident predictions than in the other cases, making the evaluation much more disadvantageous. This phenomenon is only due to the metric itself and does not appropriately reflect how hard the prediction task is: predicting short range contacts is actually easier to do than predicting long range ones.

5.2.1 CASP11 targets

Method	CASP11			CAMEO			Membrane		
	Short	Medium	Long	Short	Medium	Long	Short	Medium	Long
Proposed method	0.26	0.30	0.39	0.22	0.26	0.28	0.14	0.18	0.32
RaptorX-Contact	0.28	0.35	0.55	0.23	0.28	0.42	0.16	0.22	0.47
PconsC3	0.25	0.29	0.40	0.21	0.23	0.27	0.15	0.19	0.33
MetaPSICOV	0.26	0.31	0.39	0.22	0.22	0.28	0.16	0.21	0.35
PlmDCA	0.14	0.16	0.27	0.11	0.13	0.19	0.08	0.11	0.21
PSICOV	0.14	0.15	0.24	0.13	0.14	0.18	0.09	0.11	0.20
mfDCA	0.13	0.15	0.22	0.10	0.11	0.15	0.09	0.12	0.24

Table 5.2: Best-L PPV of different methods on short, medium and long-range contacts. Results are shown for the three different benchmark sets: CASP11 targets, CAMEO proteins, and the benchmark set of membrane proteins.

Unsupervised methods (namely DCA and PSICOV) are clearly outperformed on a very large margin by other methods. Proposed model seems to perform similarly to PConsC3, while being still significantly far from the state-of-the-art RaptorX-Contact, regardless of the evaluation metric. Tables 5.3, 5.4 and 5.5 show the overall performance on the three different test sets individually.

5.2.2 CAMEO targets

Method	Short range				Medium range				Long range			
	L/10	L/5	L/2	L	L/10	L/5	L/2	L	L/10	L/5	L/2	L
EVfold	0.25	0.21	0.15	0.12	0.33	0.27	0.19	0.13	0.37	0.33	0.25	0.19
PSICOV	0.29	0.23	0.15	0.12	0.34	0.27	0.18	0.13	0.38	0.33	0.25	0.19
plmDCA	0.35	0.28	0.17	0.12	0.40	0.32	0.21	0.14	0.43	0.39	0.31	0.23
Gremlin	0.32	0.26	0.17	0.12	0.39	0.31	0.21	0.14	0.42	0.38	0.30	0.23
CCMpred	0.35	0.27	0.17	0.12	0.40	0.31	0.21	0.14	0.44	0.40	0.31	0.23
MetaPSICOV	0.69	0.58	0.39	0.25	0.69	0.59	0.42	0.28	0.60	0.54	0.45	0.35
RaptorX-Contact	0.82	0.70	0.46	0.28	0.85	0.76	0.55	0.35	0.81	0.77	0.68	0.55
Proposed method	0.68	0.57	0.40	0.26	0.71	0.61	0.44	0.32	0.69	0.62	0.50	0.39

Table 5.3: Best-L/k PPV of different methods on the CASP11 targets with different values for L.

PlmDCA, Gremlin and CCMpred perform similarly, which is consistent with the fact that they all three rely on pseudo-likelihood direct coupling analysis. Proposed method outperforms METAPSICOV mostly on medium-range and long-range contacts on the CASP11 targets, with a 4 % increase of the best-L long-range PPV.

5.2.3 Membrane proteins

Method	Short range				Medium range				Long range			
	L/10	L/5	L/2	L	L/10	L/5	L/2	L	L/10	L/5	L/2	L
EVfold	0.17	0.13	0.11	0.09	0.23	0.19	0.13	0.10	0.25	0.22	0.17	0.13
PSICOV	0.20	0.15	0.11	0.08	0.24	0.19	0.13	0.09	0.25	0.23	0.18	0.13
plmDCA	0.22	0.16	0.11	0.09	0.27	0.22	0.14	0.10	0.30	0.26	0.20	0.15
Gremlin	0.23	0.18	0.12	0.09	0.27	0.22	0.14	0.10	0.30	0.26	0.20	0.15
CCMpred	0.21	0.17	0.11	0.08	0.27	0.22	0.14	0.10	0.31	0.26	0.20	0.15
MetaPSICOV	0.56	0.47	0.31	0.20	0.53	0.45	0.32	0.22	0.47	0.42	0.33	0.25
RaptorX-Contact	0.67	0.57	0.37	0.23	0.69	0.61	0.42	0.28	0.69	0.65	0.55	0.42
Proposed method	0.58	0.47	0.32	0.22	0.56	0.47	0.35	0.26	0.51	0.47	0.37	0.28

Table 5.4: Best-L/k PPV of different methods on the CAMEO targets with different values for L.

Proposed method outperforms METAPSICOV on the CAMEO targets, no matter whether predicted contacts are short, medium or long-range. Methods based on pseudo-likelihood DCA still perform similarly, while being more accurate than PSICOV.

Method	Short range				Medium range				Long range			
	L/10	L/5	L/2	L	L/10	L/5	L/2	L	L/10	L/5	L/2	L
EVfold	0.16	0.13	0.09	0.07	0.28	0.22	0.13	0.09	0.44	0.37	0.26	0.18
PSICOV	0.22	0.16	0.10	0.07	0.29	0.21	0.13	0.09	0.42	0.34	0.23	0.16
plmDCA	0.27	0.19	0.11	0.08	0.36	0.26	0.15	0.10	0.52	0.45	0.31	0.21
Gremlin	0.26	0.18	0.11	0.08	0.35	0.25	0.14	0.09	0.51	0.42	0.29	0.20
CCMpred	0.27	0.19	0.11	0.07	0.37	0.26	0.15	0.10	0.52	0.45	0.32	0.21
MetaPSICOV	0.45	0.35	0.22	0.14	0.49	0.40	0.27	0.18	0.61	0.55	0.42	0.30
RaptorX-Contact	0.60	0.46	0.27	0.16	0.66	0.53	0.33	0.22	0.78	0.73	0.62	0.47
Proposed method	0.38	0.30	0.21	0.14	0.45	0.37	0.26	0.18	0.58	0.52	0.41	0.32

Table 5.5: Best-L/k PPV of different methods on 398 membrane proteins with different values for L.

As can be observed in table 5.5, proposed model is comparatively less accurate when only top $L/10$ predictions are considered and gets comparatively better when k decreases. As a consequence, it performs similarly to METAPSICOV when L contacts are used. Overall, proposed model performs not as good as for the two previous datasets.

5.3 Sensitivity to the number of sequence homologs

5.3.1 CASP11 targets

Figure 5.4 shows the comparative robustness of the method to the number of sequence homologs on CASP11 targets. On the basis of the regressions, it can be observed that the method is comparatively more accurate than plmDCA on proteins with fewer sequence homologs. The phenomenon is even more present with short-range contacts, where the method performance is almost invariant to the number of homologs.

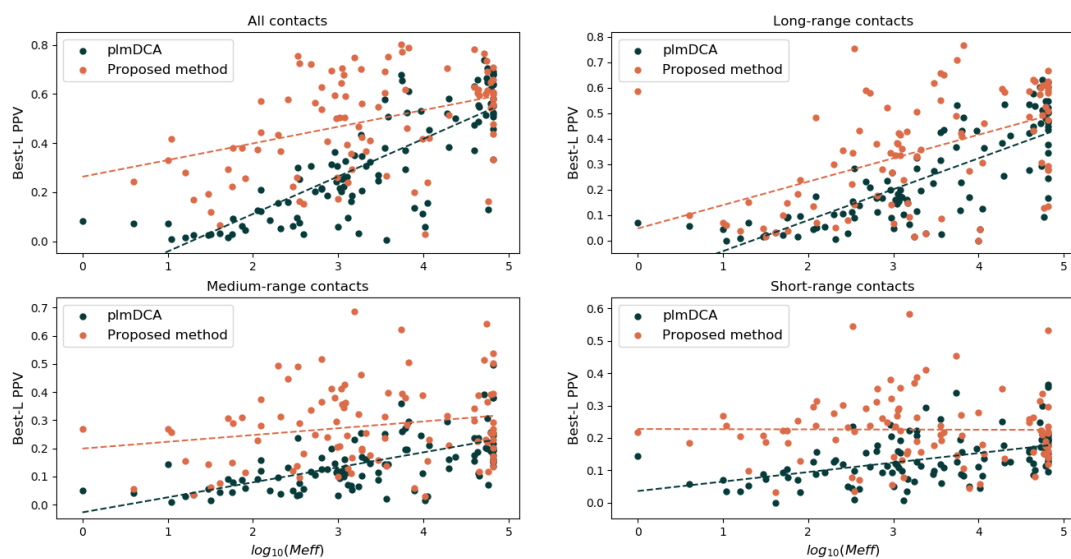


Figure 5.4: Performance as a function of the logarithm of the effective number of homologous sequences. Top figure shows the results on CASP11 targets for different types of contacts. Bottom left and bottom right figures focus on medium-range and long-range contacts, respectively.

5.3.2 CAMEO targets

When evaluated on CAMEO targets, proposed model has a significantly better overall performance compared to plmDCA, but comparable dependence on the number of sequence homologs. However it can be observed as in the case of CASP11 targets that accuracy is more robust when computed on the basis of short-range contacts.

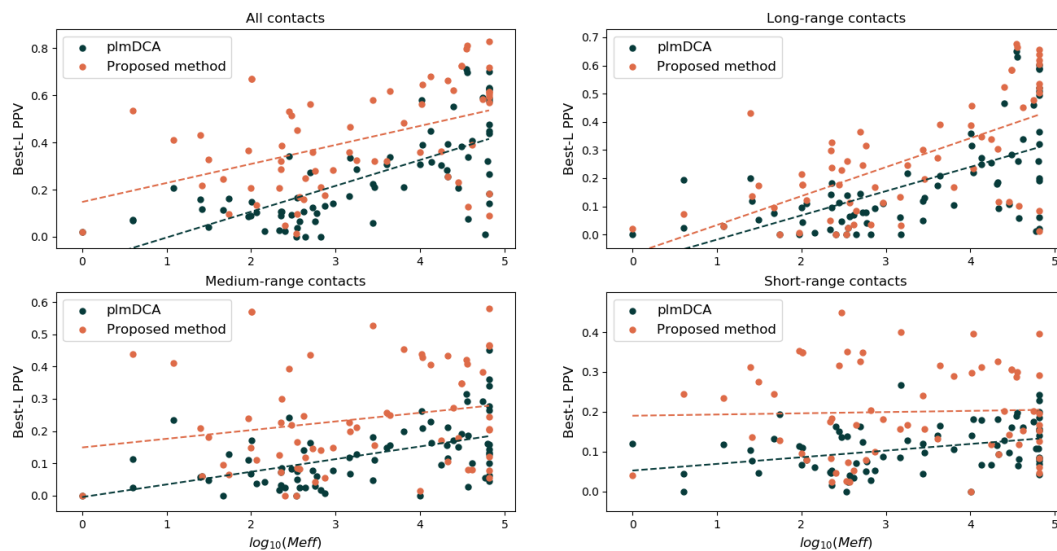


Figure 5.5: Performance as a function of the logarithm of the effective number of homologous sequences. Top figure shows the results on CAMEO targets for different types of contacts. Bottom left and bottom right figures focus on medium-range and long-range contacts, respectively.

5.3.3 Membrane proteins

Figure 5.6 confirms the observation made in the previous section: Performance gain introduced by the model w.r.t. DCA methods on membrane proteins is less important than on the two other datasets.

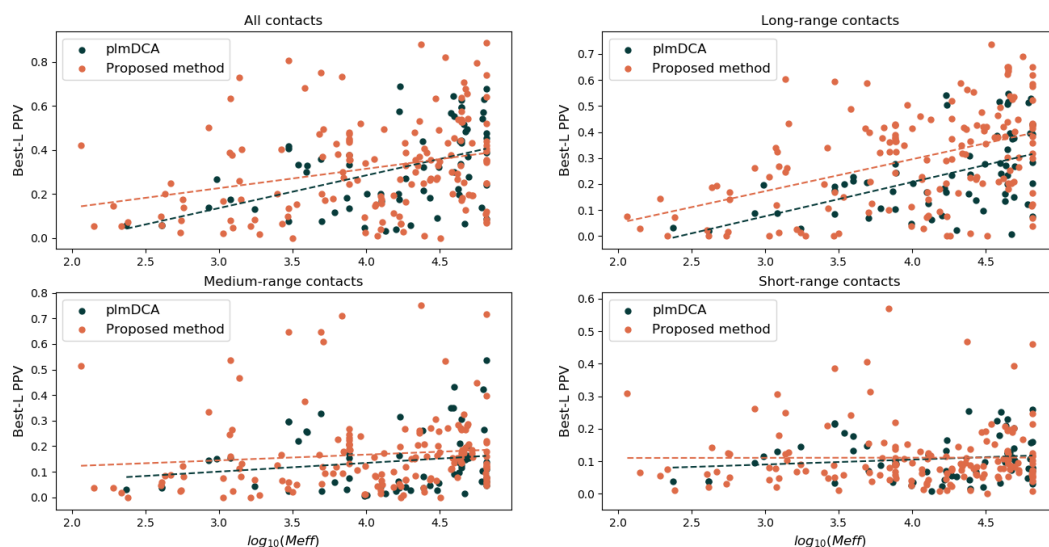


Figure 5.6: Performance as a function of the logarithm of the effective number of homologous sequences. Top figure shows the results on membrane proteins for different types of contacts. Bottom left and bottom right figures focus on medium-range and long-range contacts, respectively.

5.4 Further analysis on hard CAMEO targets

As can be observed in figure 5.7, proposed method is able to refine GaussDCA predictions from a best-L PPV of 47,2 % up to 83 % for the di-domain ARO/CYC with PDB code 4XRWA [13]. A very large number of sequence homologs are available for this protein (65535), which partly explains the success in predicting long-range contacts. The two domains of the protein lay back-to-back with their beta sheets being antiparallel. Antiparallel beta sheets are accurately predicted, but contacts between beta strands and helices do not seem to account in the top L predictions.

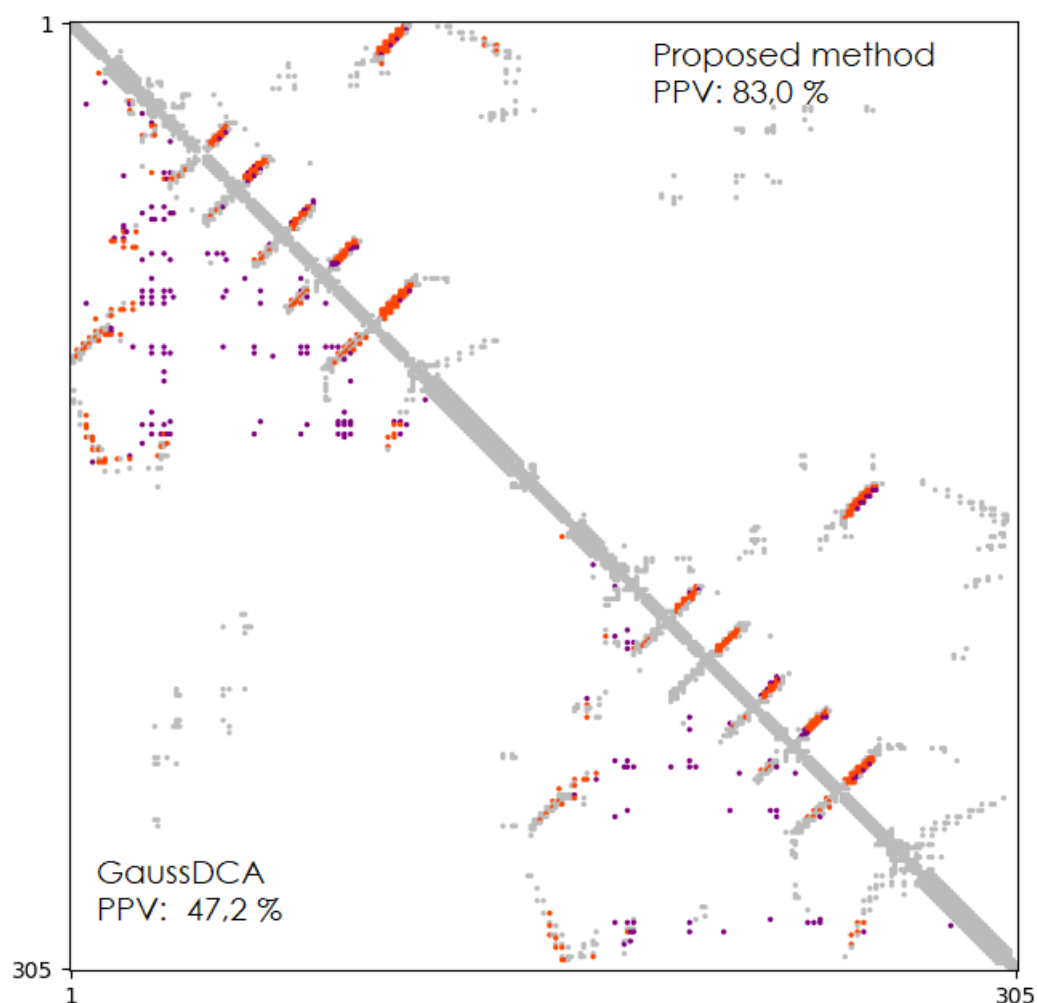


Figure 5.7: Comparison of the contact maps predicted for protein PDB:4XRWA by the proposed method and GaussDCA, respectively. True positives are highlighted in orange and false positives in purple.

It is noteworthy that the proposed model often completely fails at providing high-quality contact maps for proteins that are mainly alpha. Figure 5.8 shows predictions for the C-terminal domain of regnase-1 [104], composed of 3 α -helices. False positives are mostly short-range predicted contacts, and the model fails at finding any true contact between the first and the second helices.

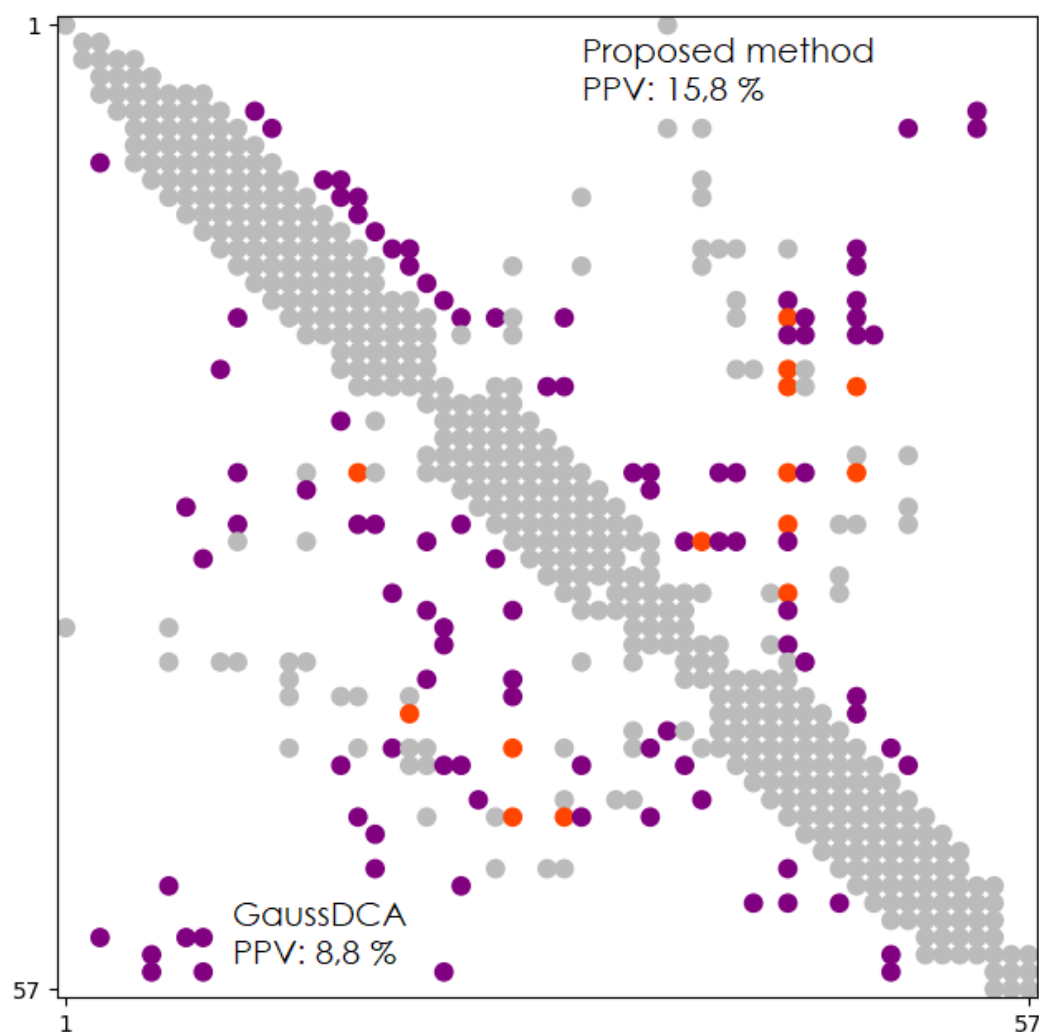


Figure 5.8: Comparison of the contact maps predicted for protein PDB:2N5LA by the proposed method and GaussDCA, respectively. True positives are highlighted in orange and false positives in purple.

Figure 5.9 is interesting as it shows predictions for a protein with 926 homologs but where most of them are much shorter. This has the effect to fill the multiple alignment matrix with gaps and yield poor evolutionary coupling predictions. Despite the ability of the model to improve contacts predicted by GaussDCA between the 120 first residues of the sequence, it remains unable to accurately predict contacts on the second part of the sequence, highlighting the limits of the refinement abilities of deep architectures.

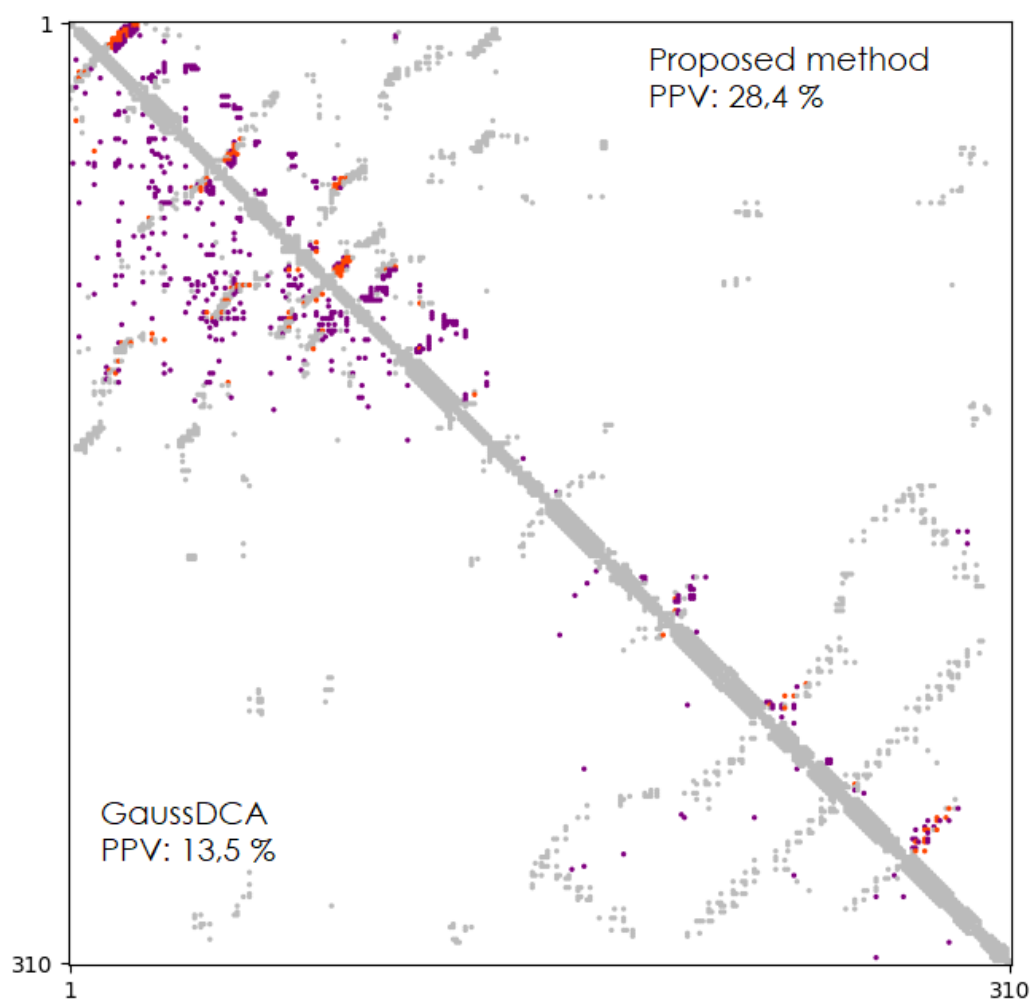


Figure 5.9: Comparison of the contact maps predicted for protein PDB:3X27D by the proposed method and GaussDCA, respectively. True positives are highlighted in orange and false positives in purple.

Chapter 6

Conclusion

The objective of this thesis has been the design of a deep neural architecture for protein contact prediction. Proposed architecture takes ideas from both RaptorX-Contact and PConsC4: it relies on a deep residual convolutional network, while still being able to predict multiple contact maps at different distance thresholds. Input features are based on target sequences, sequence homology and evolutionary couplings, in a very similar way to PConsC4. This enables the simplification of the prediction pipeline since GaussDCA (the only ECA-related predictor used in PConsC4) provides accurate DCA features in a pleasingly short amount of time.

Performance has been shown to be similar to PConsC3, while being still significantly far from the state-of-the-art RaptorX-Contact. Protein contact prediction is a very competitive area of research, and despite the fact that state of the art is not outperformed in the present work, the latter enables experimentation for future work and can be reused building more innovative approaches.

Immediate future work would focus on training the model on a larger dataset with improved preprocessing. Indeed, despite the marginal robustness of the model to the number of effective sequences in comparison with DCA methods, there is still room for additional performance that can be gained by computing statistics on much larger multiple sequence alignments. It must be noted that most of the multiple sequence alignments of training proteins come from the RaptorX-Property server, which provides a limited number of sequence homologs per target (typically < 1000). Also, now that a lot of new tertiary structure have been released after CASP12, it becomes conceivable to create a complete dataset of more than 16 000 structures.

Finally, there is still room for research now that new deep learning techniques have been developped. For example, a new unsupervised approach based on generative adversarial networks would solve the bottleneck issues related to the quality of the evolutionary coupling analysis. Generative adversarial networks have already been considered in the framework of dihedral angle prediction [51] for example. Also, the design of new end-to-end models [4] for direct structure prediction would help bypassing the two-stage prediction methods that are currently in use on all structure prediction servers, and enable the evolution of structure prediction methodology towards much simpler workflows.

Appendices

Appendix A

Structure modelling

As additional objective to this thesis, I designed a minimalist contact-assisted folding method available at: <https://github.com/AntoinePassemiers/GDE-GaussFold>.

A.1 Protein 3D alignment and evaluation metrics

$$\text{TM-score}(X^{(target)}, X^{(aligned)}) = \max_P \left[\frac{1}{L} \sum_{i=1}^L \frac{1}{1 + \left(\frac{\delta(x_i^{(target)}, P(x_i^{(aligned)}))}{\delta_0} \right)^2} \right] \quad (\text{A.1})$$

where $\delta_0 = 1.24\sqrt[3]{L - 15} - 1.8$, $\delta(x_i, y_i)$ is the euclidean distance between residue coordinates x_i and y_i , and P is a projection that preserves

The best alignment in 3D is found by determining the projection of $X^{(aligned)}$ that either maximizes the TM-score or minimizes the RMSD. Such a projection has 9 parameters:

- 3 boolean parameters that indicate whether to swap coordinates along the X, Y and Z dimensions, respectively.
- 3 real-valued parameters for translating coordinates along the X, Y and Z dimensions, respectively.
- 3 angles that parametrize the rotation matrices around the X, Y and Z axes, respectively.

$$\begin{aligned} P(x) &= R_\phi^X R_\psi^Y R_\theta^Z x + b \\ &= \begin{pmatrix} 1 & 0 & 0 \\ 0 & \cos \phi & -\sin \phi \\ 0 & \sin \phi & \cos \phi \end{pmatrix} \begin{pmatrix} \cos \psi & 0 & \sin \psi \\ 0 & 1 & 0 \\ -\sin \psi & 0 & \cos \psi \end{pmatrix} \begin{pmatrix} \cos \theta & -\sin \theta & 0 \\ \sin \theta & \cos \theta & 0 \\ 0 & 0 & 1 \end{pmatrix} x + \begin{pmatrix} b^X \\ b^Y \\ b^Z \end{pmatrix} \end{aligned}$$

A.2 Contact-assisted 3D modelling

This appendix describes the algorithm used to reconstruct proteins in three dimensions. Like GDFuzz3D [76], it uses graph distances to convert predicted contact maps in order to approximate distance maps. However, the proposed method is template-free and does not make use of MODELLER [28] as in GDFuzz3D.

A.2.1 Graph distances

Let's use the graph representation of contact maps as in section 2.1.2 about Protein Contact Networks. The graph distance between two residues is defined as the length of the shortest path between them. Predicted contact maps are converted to binary adjacency matrices by keeping only the $4.5L$ top predicted contacts.

A.2.2 Approximate Euclidean distances

As explained in [76], there is a linear relationship between graph distances and real euclidean distances. Let $GD_{i,j}$ be the graph distance between residues i and j . Then the euclidean distance $d(\mathbf{x}_i, \mathbf{x}_j)$ between the corresponding points x_i and x_j is approximated by:

$$d(\mathbf{x}_i, \mathbf{x}_j) = 5.72 \times GD_{i,j} \quad (\text{A.2})$$

It must be noted that the error on the euclidean distance also increases with graph distance. For all residue pairs for which no tractable information is available and no protein template is available, this estimated distance remains the best estimator.

A.2.3 Gaussian restraints

Each residue pair may be associated to at most one Gaussian restraint. A Gaussian restraint is defined by its mean and standard deviation, computed empirically over a set of distances with specific properties like graph distances, sequence separation or secondary structure.

As an example, residue pairs with a graph distance of one (residues are in contact) and a sequence separation of one have an average distance equal to the $C_\alpha - C_\beta$ distance (3.82 Å), and a standard deviation equal to 0.35 Å.

Restraint type	Seq. sep.	Mean	Standard deviation
Repulsion	≥ 6	20.00	120.00
Interior	-	5.00	10.00
Adjacent	1	3.81	0.1
Next adjacent	2	5.20	0.55
Next adjacent	3	7.00	0.71
Intra-alpha	1	3.82	0.35
Intra-alpha	2	5.50	0.52
Intra-alpha	3	5.33	0.93
Intra-alpha	4	6.42	1.04
Intra-beta	1	3.80	0.28
Intra-beta	2	6.66	0.30
Alpha/beta	≥ 4	6.05	0.95
Helix/coil	≥ 4	6.60	0.92
All	≥ 4	$5.72 \times \text{GD}$	$1.34 \times \text{GD}$

Table A.1: Gaussian restraints present in the 3D model

The set of points X that best satisfies Gaussian restraints is simply obtained by log-likelihood maximization:

$$\hat{X} = \underset{X}{\operatorname{argmax}} \sum_{i < j, (\mu_{i,j}, \sigma_{i,j}) \in R} \left(\frac{\delta(x_i, x_j) - \mu_{i,j}}{\sigma_{i,j}} \right)^2 \quad (\text{A.3})$$

where R is the set of parameters of the Gaussian restraints. This notation is used because all residue pairs may not be restrained.

A.2.4 Optimization algorithm

A.2.4.1 L-BFGS

L-BFGS has been used for optimizing the objective function. However, it had been observed that it is prone to getting stuck in local minima. The study of Chelvanayagam et al. [14] suggests that solving the problem in a high dimensional space solves the issue (e.g. 5 dimensions). However, this suggests that a suitable method for projecting the solution back to the 3D space is available. Instead an heuristic algorithm has been used for global convergence, coupled with L-BFGS to also ensure local convergence.

A.2.5 Evolutionary algorithm

Gaussian log-likelihood is maximized by a vanilla genetic algorithm. The initial population is obtained by adding random noise to the coordinates predicted by the multidimensional scaling algorithm. Parent selection is done by creating two random partitions from current population and keeping the individuals that maximize log-likelihood in each one. A new individual is then created by taking each point from either the first or the second parent, randomly. Mutation is simulated by adding Gaussian noise to each

point with a 50% chance. Finally, the individual with lowest log-likelihood is replaced by the newly created individual.

The set of hyper-parameters is composed of:

- Population size (default value is 2000)
- Partition size (default value is 50)
- Maximum number of iterations (default value is 200000)
- Standard deviation of mutation noise (default value is 10)

A.3 Evaluation of GDE-GaussFold

Performance on the PSICOV150 dataset is shown in figure A.1. Proposed method significantly outperforms FT-COMAR, while still being far from state-of-the-art methods. This is largely explained by the fact that the method is template-free, and that high-quality templates are available for many of the PSICOV150 families.

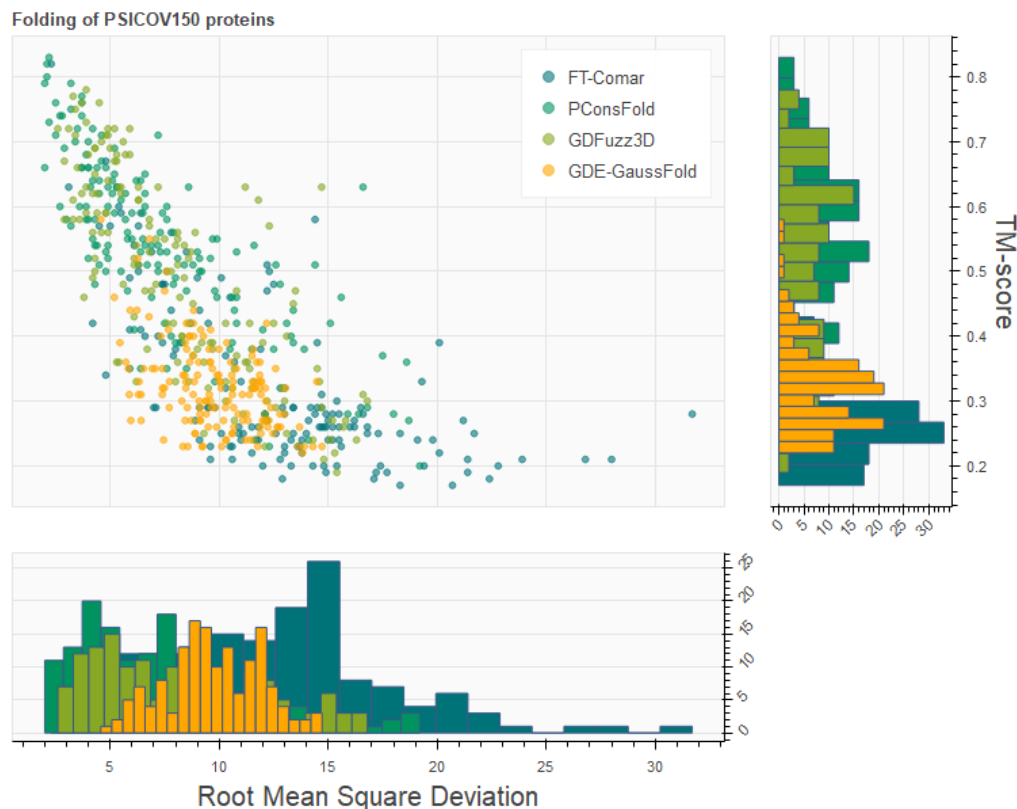


Figure A.1: TM-scores and RMSD of different folding methods including GDE-GaussFold, with density functions on the sides.

Bibliography

- [1] Badri Adhikari, Debswapna Bhattacharya, Renzhi Cao, and Jianlin Cheng. Confold: residue-residue contact-guided ab initio protein folding. *Proteins: Structure, Function, and Bioinformatics*, 83(8):1436–1449, 2015.
- [2] Badri Adhikari and Jianlin Cheng. Confold2: improved contact-driven ab initio protein structure modeling. *BMC bioinformatics*, 19(1):22, 2018.
- [3] Badri Adhikari, Jie Hou, and Jianlin Cheng. Dncon2: Improved protein contact prediction using two-level deep convolutional neural networks. 12 2017.
- [4] Mohammed AlQuraishi. End-to-end differentiable learning of protein structure. *Cell systems*, 8(4):292–301, 2019.
- [5] Christian B Anfinsen. Principles that govern the folding of protein chains. *Science*, 181(4096):223–230, 1973.
- [6] William R. Atchley, Jieping Zhao, Andrew D. Fernandes, and Tanja Drüke. Solving the protein sequence metric problem. *Proc Natl Acad Sci U S A*, 102(18):6395–6400, May 2005. 15851683[pmid].
- [7] Carlo Baldassi, Marco Zamparo, Christoph Feinauer, Andrea Procaccini, Riccardo Zecchina, Martin Weigt, and Andrea Pagnani. Fast and accurate multivariate gaussian modeling of protein families: Predicting residue contacts and protein-interaction partners. *PLOS ONE*, 9(3):1–12, 03 2014.
- [8] S. Behnel, R. Bradshaw, C. Citro, L. Dalcin, D.S. Seljebotn, and K. Smith. Cython: The best of both worlds. *Computing in Science Engineering*, 13(2):31–39, 2011.
- [9] Yoshua Bengio and Yann Lecun. *Scaling learning algorithms towards AI*. MIT Press, 2007.
- [10] James Bergstra, Brent Komer, Chris Eliasmith, Dan Yamins, and David D Cox. Hyperopt: a python library for model selection and hyperparameter optimization. *Computational Science & Discovery*, 8(1):014008, jul 2015.
- [11] James S Bergstra, Rémi Bardenet, Yoshua Bengio, and Balázs Kégl. Algorithms for hyper-parameter optimization. In *Advances in neural information processing systems*, pages 2546–2554, 2011.
- [12] R. Bollapragada, D. Mudigere, J. Nocedal, H.-J. M. Shi, and P. T. P. Tang. A Progressive Batching L-BFGS Method for Machine Learning. *ArXiv e-prints*, February 2018.

- [13] Grace Caldara-Festin, David R Jackson, Jesus F Barajas, Timothy R Valentic, Avinash B Patel, Stephanie Aguilar, MyChi Nguyen, Michael Vo, Avinash Khanna, Eita Sasaki, et al. Structural and functional analysis of two di-domain aromatase/cyclases from type ii polyketide synthases. *Proceedings of the National Academy of Sciences*, 112(50):E6844–E6851, 2015.
- [14] Gareth Chelvanayagam, Lukas Knecht, Thomas Jenny, Steven A Benner, and Gaston H Gonnet. A combinatorial distance-constraint approach to predicting protein tertiary models from known secondary structure. *Folding and Design*, 3(3):149–160, 1998.
- [15] Hua Cheng, R. Dustin Schaeffer, Yuxing Liao, Lisa N. Kinch, Jimin Pei, Shuoyong Shi, Bong-Hyun Kim, and Nick V. Grishin. Ecod: An evolutionary classification of protein domains. *PLOS Computational Biology*, 10(12):1–18, 12 2014.
- [16] Axel Cleeremans, David Servan-Schreiber, and James L McClelland. Finite state automata and simple recurrent networks. *Neural computation*, 1(3):372–381, 1989.
- [17] R. Collobert, K. Kavukcuoglu, and C. Farabet. Torch7: A matlab-like environment for machine learning. In *BigLearn, NIPS Workshop*, 2011.
- [18] Pietro Di Lena, Ken Nagata, and Pierre Baldi. Deep architectures for protein contact map prediction. *Bioinformatics*, 28(19):2449–2457, 2012.
- [19] L. Di Paola, M. De Ruvo, P. Paci, D. Santoni, and A. Giuliani. Protein contact networks: An emerging paradigm in chemistry. *Chemical Reviews*, 113(3):1598–1613, 2013. PMID: 23186336.
- [20] Wenze Ding, Wenzhi Mao, Di Shao, Wenxuan Zhang, and Haipeng Gong. Deep-conpred2: An improved method for the prediction of protein residue contacts. *Computational and Structural Biotechnology Journal*, 16, 11 2018.
- [21] John Duchi, Elad Hazan, and Yoram Singer. Adaptive subgradient methods for on-line learning and stochastic optimization. *Journal of Machine Learning Research*, 12(Jul):2121–2159, 2011.
- [22] Stanley D Dunn, Lindi M Wahl, and Gregory B Gloor. Mutual information without the influence of phylogeny or entropy dramatically improves residue contact prediction. *Bioinformatics*, 24(3):333–340, 2007.
- [23] Sean R. Eddy. Profile hidden markov models. *Bioinformatics (Oxford, England)*, 14(9):755–763, 1998.
- [24] Magnus Ekeberg, Tuomo Hartonen, and Erik Aurell. Fast pseudolikelihood maximization for direct-coupling analysis of protein structure from many homologous amino-acid sequences. *Journal of Computational Physics*, 276:341 – 356, 2014.
- [25] Magnus Ekeberg, Cecilia Lövkvist, Yueheng Lan, Martin Weigt, and Erik Aurell. Improved contact prediction in proteins: Using pseudolikelihoods to infer potts models. *Phys. Rev. E*, 87:012707, Jan 2013.
- [26] Sara El-Gebali, Jaina Mistry, Alex Bateman, Sean R Eddy, Aurélien Luciani, Simon C Potter, Matloob Qureshi, Lorna J Richardson, Gustavo A Salazar, Alfredo

- Smart, Erik L L Sonhammer, Layla Hirsh, Lisanna Paladin, Damiano Piovesan, Silvio C E Tosatto, and Robert D Finn. The pfam protein families database in 2019. *Nucleic Acids Research*, page gky995, 2018.
- [27] R John Ellis and Saskia M Van der Vies. Molecular chaperones. *Annual review of biochemistry*, 60(1):321–347, 1991.
- [28] Narayanan Eswar, Ben Webb, Marc A Marti-Renom, MS Madhusudhan, David Eramian, Min-yi Shen, Ursula Pieper, and Andrej Sali. Comparative protein structure modeling using modeller. *Current protocols in bioinformatics*, 15(1):5–6, 2006.
- [29] R. Evans, J. Jumper, J. Kirkpatrick, L. Sifre, T.F.G. Green, C. Qin, A. Zidek, A. Nelson, A. Bridgland, H. Penedones, S. Petersen, K. Simonyan, S. Crossan, D.T. Jones, D. Silver, K. Kavukcuoglu, D. Hassabis, and A.W. Senior. De novo structure prediction with deep-learning based scoring.
- [30] Iakes Ezkurdia, Osvaldo Grana, José MG Izarzugaza, and Michael L Tress. Assessment of domain boundary predictions and the prediction of intramolecular contacts in casp8. *Proteins: Structure, Function, and Bioinformatics*, 77(S9):196–209, 2009.
- [31] Gianluigi Forloni, Liana Terreni, Ilaria Bertani, Sergio Fogliarino, Roberto Invernizzi, Andrea Assini, Giuseppe Ribizzi, Alessandro Negro, Elena Calabrese, Maria Antonietta Volonté, et al. Protein misfolding in alzheimer’s and parkinson’s disease: genetics and molecular mechanisms. *Neurobiology of aging*, 23(5):957–976, 2002.
- [32] Jerome Friedman, Trevor Hastie, and Robert Tibshirani. Sparse inverse covariance estimation with the graphical lasso. *Biostatistics (Oxford, England)*, 9:432–41, 08 2008.
- [33] Alberto Garcia-Garcia, Sergio Orts-Escolano, Sergiu Oprea, Victor Villena-Martinez, and José García Rodríguez. A review on deep learning techniques applied to semantic segmentation. *CoRR*, abs/1704.06857, 2017.
- [34] Vladimir Golkov, Marcin J Skwark, Antonij Golkov, Alexey Dosovitskiy, Thomas Brox, Jens Meiler, and Daniel Cremers. Protein contact prediction from amino acid co-evolution using convolutional networks for graph-valued images. In *Advances in Neural Information Processing Systems*, pages 4222–4230, 2016.
- [35] Alex Graves. Generating sequences with recurrent neural networks. *arXiv preprint arXiv:1308.0850*, 2013.
- [36] Juergen Haas, Steven Roth, Konstantin Arnold, Florian Kiefer, Tobias Schmidt, Lorenza Bordoli, and Torsten Schwede. The protein model portal—a comprehensive resource for protein structure and model information. *Database*, 2013, 2013.
- [37] Nicholas Hamilton and Thomas Huber. *An Introduction to Protein Contact Prediction*, pages 87–104. Humana Press, Totowa, NJ, 2008.
- [38] Jack Hanson, Kuldip Paliwal, Thomas Litfin, Yuedong Yang, and Yaoqi Zhou. Accurate prediction of protein contact maps by coupling residual two-dimensional

- bidirectional long short-term memory with convolutional neural networks. *Bioinformatics*, 34(23):4039–4045, 2018.
- [39] William E Hart and Sorin Istrail. Robust proofs of np-hardness for protein folding: general lattices and energy potentials. *Journal of Computational Biology*, 4(1):1–22, 1997.
 - [40] Kaiming He, Xiangyu Zhang, Shaoqing Ren, and Jian Sun. Deep residual learning for image recognition. *CoRR*, abs/1512.03385, 2015.
 - [41] Sepp Hochreiter and Jürgen Schmidhuber. Long short-term memory. *Neural computation*, 9(8):1735–1780, 1997.
 - [42] Kurt Hornik. Approximation capabilities of multilayer feedforward networks. *Neural networks*, 4(2):251–257, 1991.
 - [43] Gao Huang, Zhuang Liu, Laurens Van Der Maaten, and Kilian Q Weinberger. Densely connected convolutional networks. In *Proceedings of the IEEE conference on computer vision and pattern recognition*, pages 4700–4708, 2017.
 - [44] D. H. Hubel and T. N. Wiesel. Receptive fields, binocular interaction and functional architecture in the cat’s visual cortex. *J Physiol*, 160(1):106–154.2, Jan 1962. 14449617[pmid].
 - [45] Sergey Ioffe and Christian Szegedy. Batch normalization: Accelerating deep network training by reducing internal covariate shift. *CoRR*, abs/1502.03167, 2015.
 - [46] David T Jones. Protein secondary structure prediction based on position-specific scoring matrices. *Journal of molecular biology*, 292(2):195–202, 1999.
 - [47] David T. Jones, Daniel W. A. Buchan, Domenico Cozzetto, and Massimiliano Pontil. Psicov: precise structural contact prediction using sparse inverse covariance estimation on large multiple sequence alignments. *Bioinformatics*, 28(2):184–190, 2012.
 - [48] David T Jones and Shaun M Kandathil. High precision in protein contact prediction using fully convolutional neural networks and minimal sequence features. *Bioinformatics*, page bty341, 2018.
 - [49] David T. Jones, Tanya Singh, Tomasz Kosciolk, and Stuart Tetchner. Metapsicov: combining coevolution methods for accurate prediction of contacts and long range hydrogen bonding in proteins. *Bioinformatics*, 31(7):999–1006, 2015.
 - [50] David E Kim, Frank DiMaio, Ray Yu-Ruei Wang, Yifan Song, and David Baker. One contact for every twelve residues allows robust and accurate topology-level protein structure modeling. *Proteins: Structure, Function, and Bioinformatics*, 82:208–218, 2014.
 - [51] Hyeonki Kim. Dihedral angle prediction using generative adversarial networks. *arXiv preprint arXiv:1803.10996*, 2018.
 - [52] Lisa N. Kinch, Wenlin Li, R. Dustin Schaeffer, Roland L. Dunbrack, Bohdan Monastyrskyy, Andriy Kryshchak, and Nick V. Grishin. Casp 11 target classification. *Proteins: Structure, Function, and Bioinformatics*, 84(S1):20–33, 2016.

- [53] Diederik P. Kingma and Jimmy Ba. Adam: A method for stochastic optimization. *CoRR*, abs/1412.6980, 2014.
- [54] Yann LeCun, Yoshua Bengio, and Geoffrey Hinton. Deep learning. *nature*, 521(7553):436, 2015.
- [55] Yann LeCun, Léon Bottou, Yoshua Bengio, and Patrick Haffner. Gradient-based learning applied to document recognition. *Proceedings of the IEEE*, 86(11):2278–2324, 1998.
- [56] Yann LeCun, D Touresky, G Hinton, and T Sejnowski. A theoretical framework for back-propagation. In *Proceedings of the 1988 connectionist models summer school*, volume 1, pages 21–28. CMU, Pittsburgh, Pa: Morgan Kaufmann, 1988.
- [57] Yann A. LeCun, Léon Bottou, Genevieve B. Orr, and Klaus-Robert Müller. *Efficient BackProp*, pages 9–48. Springer Berlin Heidelberg, Berlin, Heidelberg, 2012.
- [58] C Levinthal. How to fold graciously. mossbauer spectroscopy in biological systems: Proceedings of a meeting held at allerton house. *Monticello, Illinois (Debrunnder JTP, Munck E., eds.)*, pages 22–24, 1969.
- [59] Zhaoyu Li, Son P. Nguyen, Dong Xu, and Yi Shang. Protein loop modeling using deep generative adversarial network, 11 2017.
- [60] Yang Liu, Perry Palmedo, Qing Ye, Bonnie Berger, and Jian Peng. Enhancing evolutionary couplings with deep convolutional neural networks. *Cell Systems*, 6(1), Jan 2018.
- [61] Lorenzo Livi, Enrico Maiorino, Alessandro Giuliani, Antonello Rizzi, and Alireza Sadeghian. A generative model for protein contact networks. *Journal of Biomolecular Structure and Dynamics*, 34(7):1441–1454, 2016. PMID: 26474097.
- [62] Grzegorz Markowski, Krzysztof Grabczewski, and Rafal Adamczak. Oversampling negative class improves contact map prediction. *Int J Pharma Med Biol Sci*, 5(4):211–216, 2016.
- [63] Mirco Michel, Sikander Hayat, Marcin J Skwark, Chris Sander, Debora S Marks, and Arne Elofsson. Pconsfold: improved contact predictions improve protein models. *Bioinformatics*, 30(17):i482–i488, 2014.
- [64] Mirco Michel, David Menendez Hurtado, and Arne Elofsson. Pconsc4: fast, free, easy, and accurate contact predictions. *bioRxiv*, 2018.
- [65] Alex D Michie, Christine A Orengo, and Janet M Thornton. Analysis of domain structural class using an automated class assignment protocol. *Journal of molecular biology*, 262(2):168–185, 1996.
- [66] Faruck Morcos, Terence Hwa, José N. Onuchic, and Martin Weigt. *Direct Coupling Analysis for Protein Contact Prediction*, pages 55–70. Springer New York, New York, NY, 2014.
- [67] Faruck Morcos, Terence Hwa, José N Onuchic, and Martin Weigt. Direct coupling analysis for protein contact prediction. In *Protein structure prediction*, pages 55–70. Springer, 2014.

- [68] Faruck Morcos, Andrea Pagnani, Bryan Lunt, Arianna Bertolino, Debora S. Marks, Chris Sander, Riccardo Zecchina, José N. Onuchic, Terence Hwa, and Martin Weigt. Direct-coupling analysis of residue coevolution captures native contacts across many protein families. *Proceedings of the National Academy of Sciences*, 108(49):E1293–E1301, 2011.
- [69] John Moult, Krzysztof Fidelis, Andriy Kryshchak, Torsten Schwede, and Anna Tramontano. Critical assessment of methods of protein structure prediction (casp) — round x. *Proteins: Structure, Function, and Bioinformatics*, 82(S2):1–6, 2014.
- [70] John Moult, Krzysztof Fidelis, Andriy Kryshchak, Torsten Schwede, and Anna Tramontano. Critical assessment of methods of protein structure prediction: Progress and new directions in round xi. *Proteins: Structure, Function, and Bioinformatics*, 84(S1):4–14, 2016.
- [71] John Moult, Krzysztof Fidelis, Andriy Kryshchak, Torsten Schwede, and Anna Tramontano. Critical assessment of methods of protein structure prediction (casp)—round xii. *Proteins: Structure, Function, and Bioinformatics*, 86:7–15, 2018.
- [72] Osvaldo Olmea and Alfonso Valencia. Improving contact predictions by the combination of correlated mutations and other sources of sequence information. *Folding and Design*, 2:S25 – S32, 1997.
- [73] Christine A Orengo, Alex D Michie, Susan Jones, David T Jones, Mark B Swindells, and Janet M Thornton. Cath—a hierarchic classification of protein domain structures. *Structure*, 5(8):1093–1109, 1997.
- [74] Jian Peng and Jinbo Xu. Raptorx: exploiting structure information for protein alignment by statistical inference. *Proteins: Structure, Function, and Bioinformatics*, 79(S10):161–171, 2011.
- [75] Na Pi and Julie A Leary. Determination of enzyme/substrate specificity constants using a multiple substrate esi-ms assay. *Journal of the American Society for Mass Spectrometry*, 15(2):233–243, 2004.
- [76] Michal J Pietal, Janusz M Bujnicki, and Lukasz P Kozłowski. Gdfuzz3d: a method for protein 3d structure reconstruction from contact maps, based on a non-euclidean distance function. *Bioinformatics*, 31(21):3499–3505, 2015.
- [77] MG Reese, Ole Lund, Jakob Bohr, Henrik Bohr, JE Hansen, and Søren Brunak. Distance distributions in proteins: a six-parameter representation. *Protein Engineering, Design and Selection*, 9(9):733–740, 1996.
- [78] Michael Remmert, Andreas Biegert, Andreas Hauser, and Johannes Söding. Hhblits: lightning-fast iterative protein sequence searching by hmm-hmm alignment. *Nature Methods*, 9:173 EP –, Dec 2011.
- [79] Olaf Ronneberger, Philipp Fischer, and Thomas Brox. U-net: Convolutional networks for biomedical image segmentation. *CoRR*, abs/1505.04597, 2015.
- [80] Sebastian Ruder. An overview of gradient descent optimization algorithms. *CoRR*, abs/1609.04747, 2016.

- [81] David E Rumelhart, Geoffrey E Hinton, Ronald J Williams, et al. Learning representations by back-propagating errors. *Cognitive modeling*, 5(3):1, 1988.
- [82] Ida Schomburg, Antje Chang, Sandra Placzek, Carola Söhngen, Michael Rother, Maren Lang, Cornelia Munaretto, Susanne Ulas, Michael Stelzer, Andreas Grote, et al. Brenda in 2013: integrated reactions, kinetic data, enzyme function data, improved disease classification: new options and contents in brenda. *Nucleic acids research*, 41(D1):D764–D772, 2012.
- [83] Stefan Seemayer, Markus Gruber, and Johannes Söding. Ccmpred–fast and precise prediction of protein residue-residue contacts from correlated mutations. *Bioinformatics*, 30(21):3128–3130, Nov 2014. 25064567[pmid].
- [84] Robert Sheridan, Robert J. Fieldhouse, Sikander Hayat, Yichao Sun, Yevgeniy Antipin, Li Yang, Thomas Hopf, Debora S. Marks, and Chris Sander. Evfold.org: Evolutionary couplings and protein 3d structure prediction. *bioRxiv*, 2015.
- [85] A Shrake and JA Rupley. Environment and exposure to solvent of protein atoms. lysozyme and insulin. *Journal of molecular biology*, 79(2):351–371, 1973.
- [86] Karen Simonyan, Andrea Vedaldi, and Andrew Zisserman. Deep inside convolutional networks: Visualising image classification models and saliency maps. *CoRR*, abs/1312.6034, 2013.
- [87] Marcin J. Skwark, Abbi Abdel-Rehim, and Arne Elofsson. Pconsc: combination of direct information methods and alignments improves contact prediction. *Bioinformatics*, 29(14):1815–1816, 2013.
- [88] Marcin J Skwark, Mirco Michel, David Menendez Hurtado, Magnus Ekeberg, and Arne Elofsson. Accurate contact predictions for thousands of protein families using pconsc3. *bioRxiv*, 2016.
- [89] Marcin J. Skwark, Daniele Raimondi, Mirco Michel, and Arne Elofsson. Improved contact predictions using the recognition of protein like contact patterns. *PLOS Computational Biology*, 10(11):1–14, 11 2014.
- [90] Siqi Sun, Jianzhu Ma, Sheng Wang, and Jinbo Xu. Predicting diverse m-best protein contact maps. In *2015 IEEE International Conference on Bioinformatics and Biomedicine (BIBM)*, pages 1287–1295. IEEE, 2015.
- [91] Máttyás A Sustik and Ben Calderhead. Glassofast: an efficient glasso implementation. *UTCS Technical Report TR-12-29 2012*, 2012.
- [92] William R Taylor, Russell S Hamilton, and Michael I Sadowski. Prediction of contacts from correlated sequence substitutions. *Current Opinion in Structural Biology*, 23(3):473 – 479, 2013. New constructs and expressions of proteins / Sequences and topology.
- [93] Nikos Tsardakas Renhuldt. Protein contact prediction based on the tiramisu deep learning architecture. Master’s thesis, KTH, School of Electrical Engineering and Computer Science (EECS), 2018.

- [94] Marco Vassura, Luciano Margara, Pietro Di Lena, Filippo Medri, Piero Fariselli, and Rita Casadio. Ft-comar: fault tolerant three-dimensional structure reconstruction from protein contact maps. *Bioinformatics*, 24(10):1313–1315, 2008.
- [95] Marco Vassura, Luciano Margara, Filippo Medri, Pietro di Lena, Piero Fariselli, and Rita Casadio. Reconstruction of 3d structures from protein contact maps. In Ion Măndoiu and Alexander Zelikovsky, editors, *Bioinformatics Research and Applications*, pages 578–589, Berlin, Heidelberg, 2007. Springer Berlin Heidelberg.
- [96] Guoli Wang and Roland L Dunbrack Jr. Pisces: a protein sequence culling server. *Bioinformatics*, 19(12):1589–1591, 2003.
- [97] Sheng Wang, Wei Li, Shiwang Liu, and Jinbo Xu. Raptorx-property: a web server for protein structure property prediction. *Nucleic acids research*, 44(W1):W430–W435, 2016.
- [98] Sheng Wang, Wei Li, Renyu Zhang, Shiwang Liu, and Jinbo Xu. Coinfold: a web server for protein contact prediction and contact-assisted protein folding. *Nucleic acids research*, 44(W1):W361–W366, 2016.
- [99] Sheng Wang, Siqi Sun, Zhen Li, Renyu Zhang, and Jinbo Xu. Accurate de novo prediction of protein contact map by ultra-deep learning model. *PLOS Computational Biology*, 13(1):1–34, 01 2017.
- [100] Zhiyong Wang and Jinbo Xu. Predicting protein contact map using evolutionary and physical constraints by integer programming. *Bioinformatics*, 29(13):i266–i273, Jul 2013. 23812992[pmid].
- [101] Zhiyong Wang and Jinbo Xu. Predicting protein contact map using evolutionary and physical constraints by integer programming. *Bioinformatics*, 29(13):i266–i273, Jul 2013. btt211[PII].
- [102] Martin Weigt, Robert A. White, Hendrik Szurmant, James A. Hoch, and Terence Hwa. Identification of direct residue contacts in protein-protein interaction by message passing. *Proc Natl Acad Sci U S A*, 106(1):67–72, Jan 2009.
- [103] Dapeng Xiong, Jianyang Zeng, and Haipeng Gong. A deep learning framework for improving long-range residue–residue contact prediction using a hierarchical strategy. *Bioinformatics*, 33(17):2675–2683, 2017.
- [104] Mariko Yokogawa, Takashi Tsushima, Nobuo N Noda, Hiroyuki Kumeta, Yoshiaki Enokizono, Kazuo Yamashita, Daron M Standley, Osamu Takeuchi, Shizuo Akira, and Fuyuhiko Inagaki. Structural basis for the regulation of enzymatic activity of regnase-1 by domain-domain interactions. *Scientific reports*, 6:22324, 2016.
- [105] Matthew D. Zeiler. ADADELTA: an adaptive learning rate method. *CoRR*, abs/1212.5701, 2012.